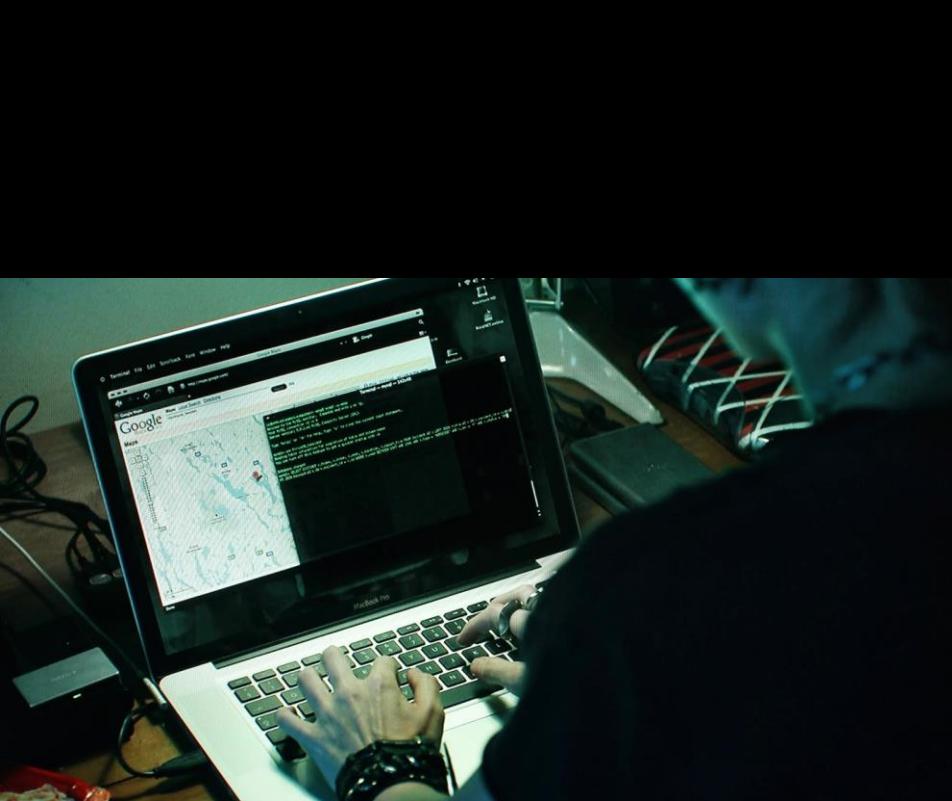


The RTOS
Exploit Mitigation
Blues

Jos Wetzels

HARDWEAR.IO



Jos Wetzels



Midnight Blue

Independent
Security Researcher

UNIVERSITY
OF TWENTE.

(Former)
Security Researcher

ICS

Critical Infrastructure

IoT

Embedded BinSec

Automotive

HIDS / NIDS

Medical

Access Controls

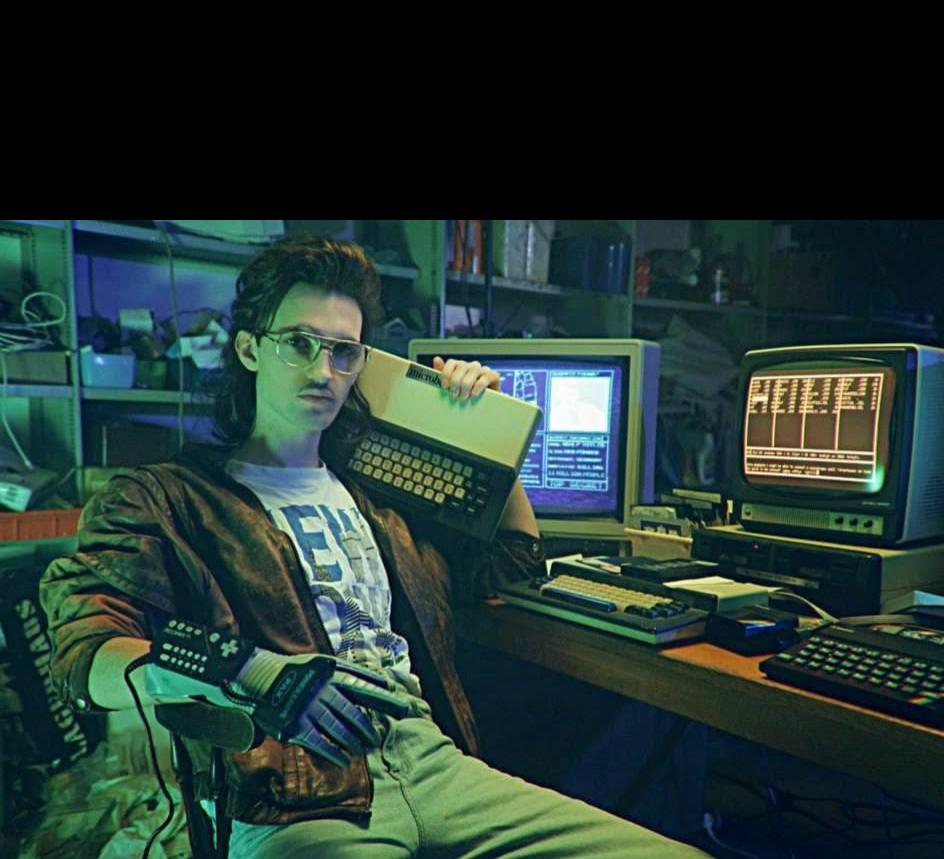
Networking & Firewalls



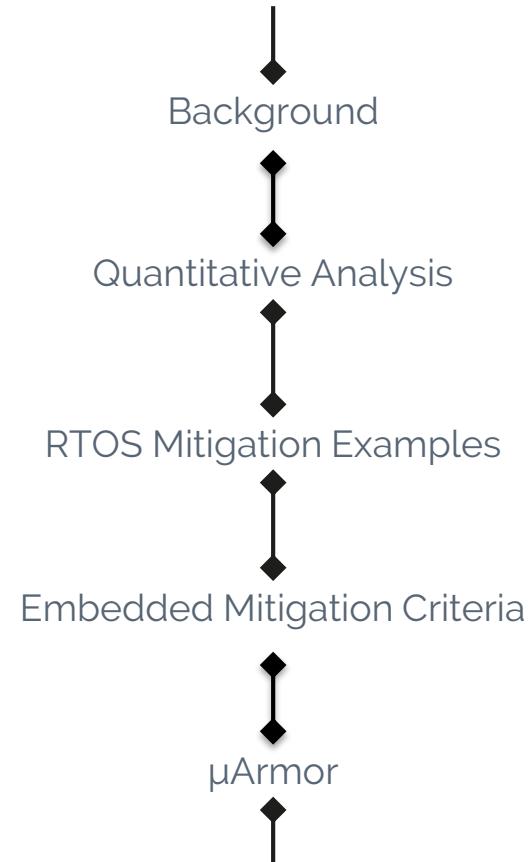
@s4mvartaka



www.midnightbluelabs.com
samvartaka.github.io



ROADMAP



What this talk is not about



What this talk is also not about

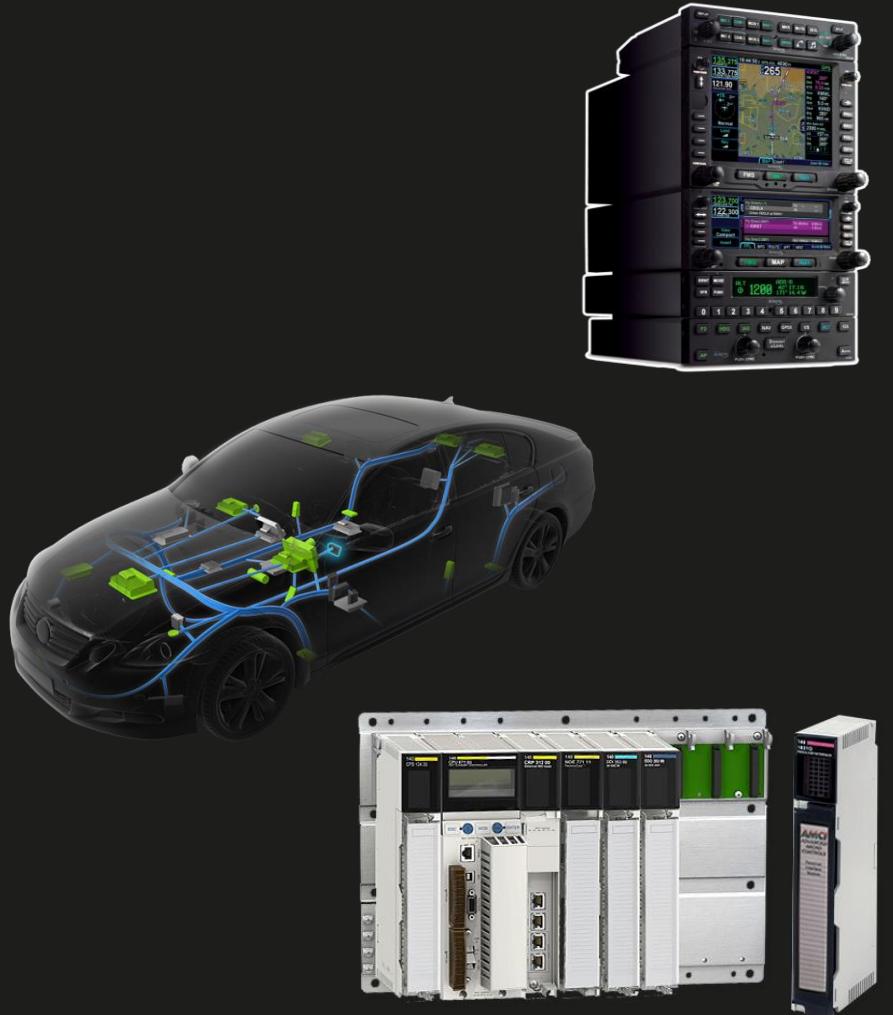


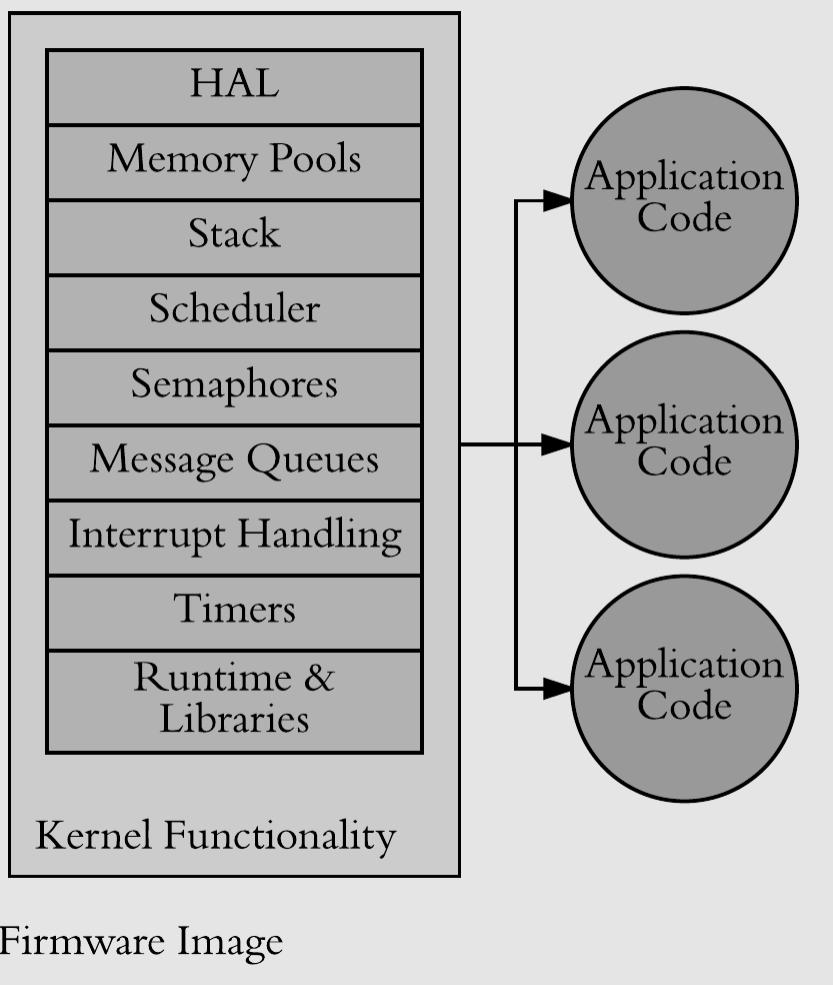
What this talk *is* about



Real-Time Operating System (RTOS)

- Guarantee Response within Time Constraints
Important in safety-critical sys.
- Predictability & Determinism > Speed
- Soft RT → Usefulness of result degrades after deadline
- Hard RT → Missing deadline = system failure



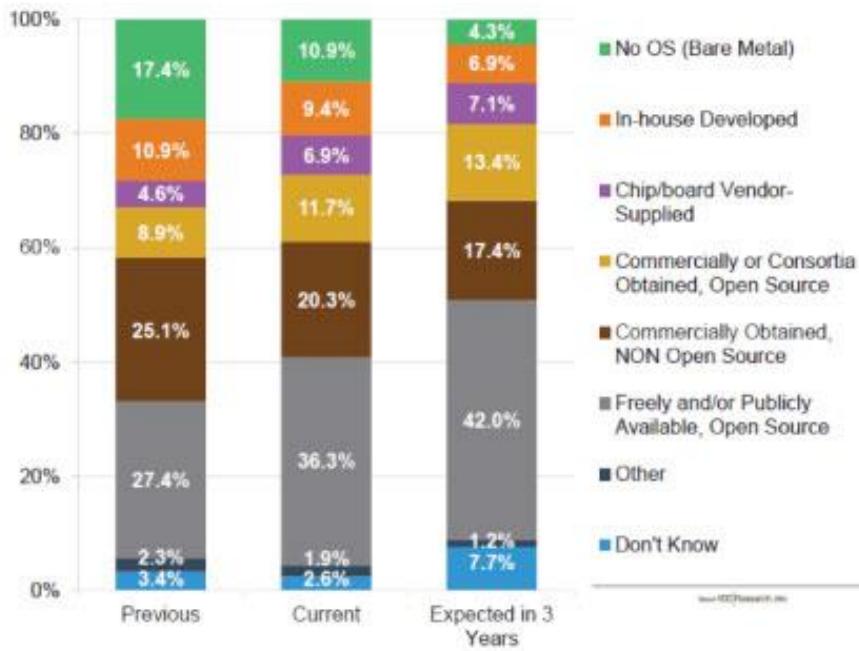


Library / Unikernel OS

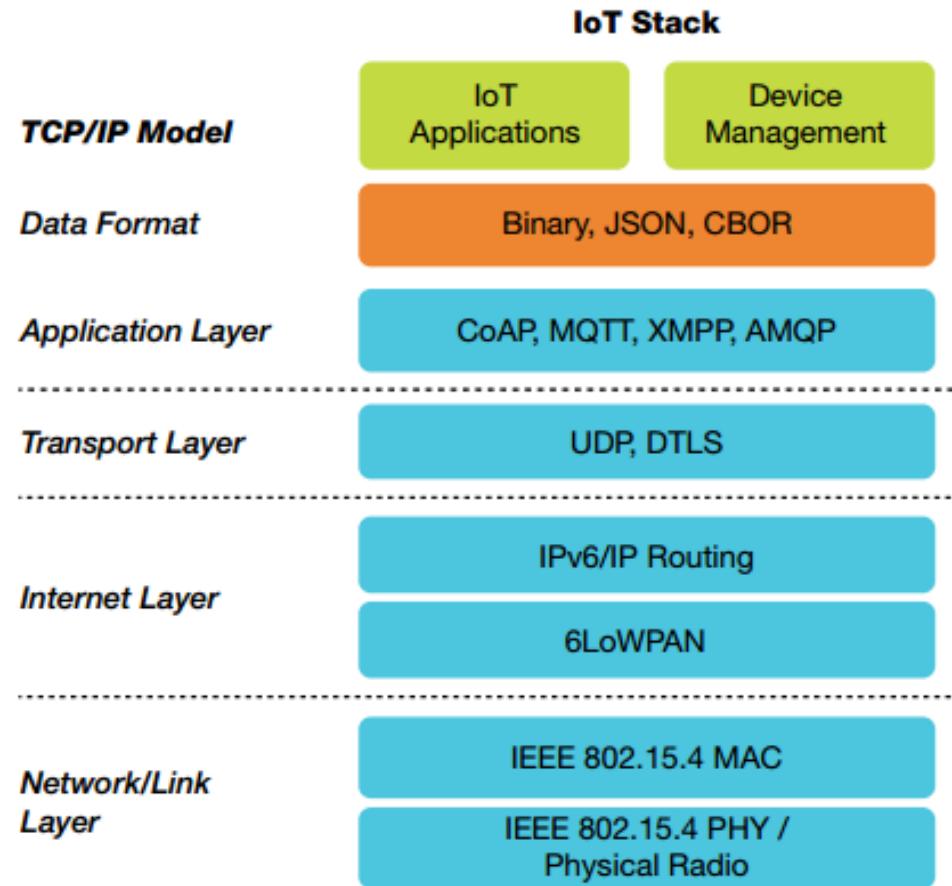
- **OS as collection of libraries**
Linked with application code into single address space image
- **No loader, Syscalls implemented as function calls**
Saves context-switching

Rising Embedded RTOS Usage

Past, Present, and Expected Primary Operating System, Segmented by Primary OS Source
(Percent of Respondents)



*Note: Interim data from VDC's 2016 Embedded Engineer Survey.



Increasing Connectivity

- Protocol Stacks Everywhere

Talk to gateway, other nodes, smartphones, integration with cloud, etc.

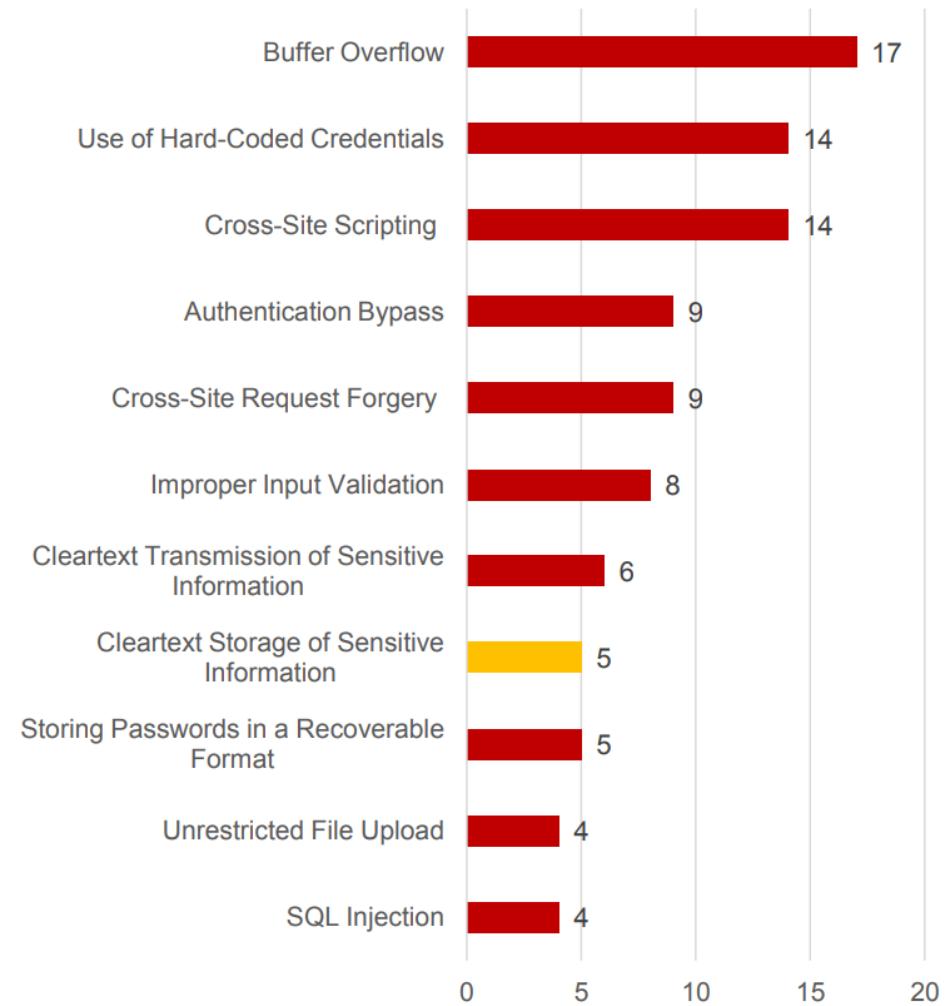
Increasing Complexity

# Processors	Bare Metal	RTOS	Linux
1	33%	18%	13%
2-3	22%	24%	17%
4+	11%	25%	32%

- More processors, memory, peripherals
- Simple bare metal control loops & state machines not enough
- Hence rise of RTOS

A photograph of two individuals, a man and a woman, working at a computer keyboard in a dimly lit environment. The man is in the foreground, looking down at the keyboard. The woman is standing behind him, also focused on the screen. The scene has a blue-tinted, cinematic feel.

Embedded Binary Security



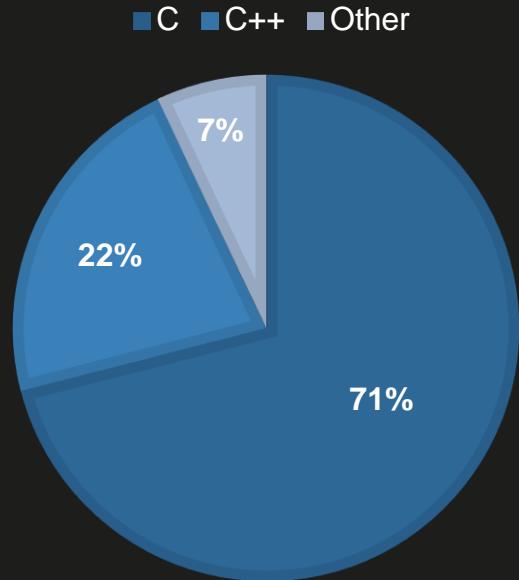
Memory
Corruption is
a Big Deal™
in Embedded



Heightened Impact

- **Patching Issues**
 - Fragmented Responsibility
 - Lack of Secure Update
 - Availability / Safety
- **Long Exposure Windows**
 - 'Forever Days'
- **No Privilege Separation**
 - Flat Memory Model
- **Exploits Scale Well**
 - Single protocol driver bug can affect billions of devices

PRIMARY EMBEDDED PROGRAMMING LANGUAGE



Unsafe Languages Are Here To Stay

- Ideally use safe languages
- But unsafe continues to dominate

Despite old (Ada) and upcoming (Rust) embedded-suited langs



The background image depicts a detailed architectural rendering of a medieval castle. It features a massive outer wall with several cylindrical towers and a central, more complex section with a rounded tower and a smaller building attached. The castle is built from light-colored stone blocks and has a series of walkways and battlements along its perimeter. A central courtyard is visible, enclosed by the castle's walls.

Exploit Mitigations

General Purpose Exploitation Has Been Getting Harder



VS



2010

2016

What About Embedded?





Right...

RTOS Exploit Examples

Name	Target	Type	RTOS	Mitigations
BroadPWN / CVE-2017- 9417	Broadcom WiFi	Heap Buffer Overflow	VxWorks	NO
CVE-2017-0561	Broadcom WiFi	Heap Buffer Overflow	VxWorks	NO
VxWorks DHCP / DNS	VxWorks DHCP / DNS Client & Server	Heap / Stack Buffer Overflow	VxWorks	NO
CVE-2015-7599	VxWorks RPC	Integer Overflow	VxWorks	NO
CVE-2010-3832	Intel/Comneon Baseband	Heap Buffer Overflow	ThreadX / Nucleus+	NO¹
Qualcomm UMTS AUTN	Qualcomm Baseband	Stack Buffer Overflow	REX	NO²

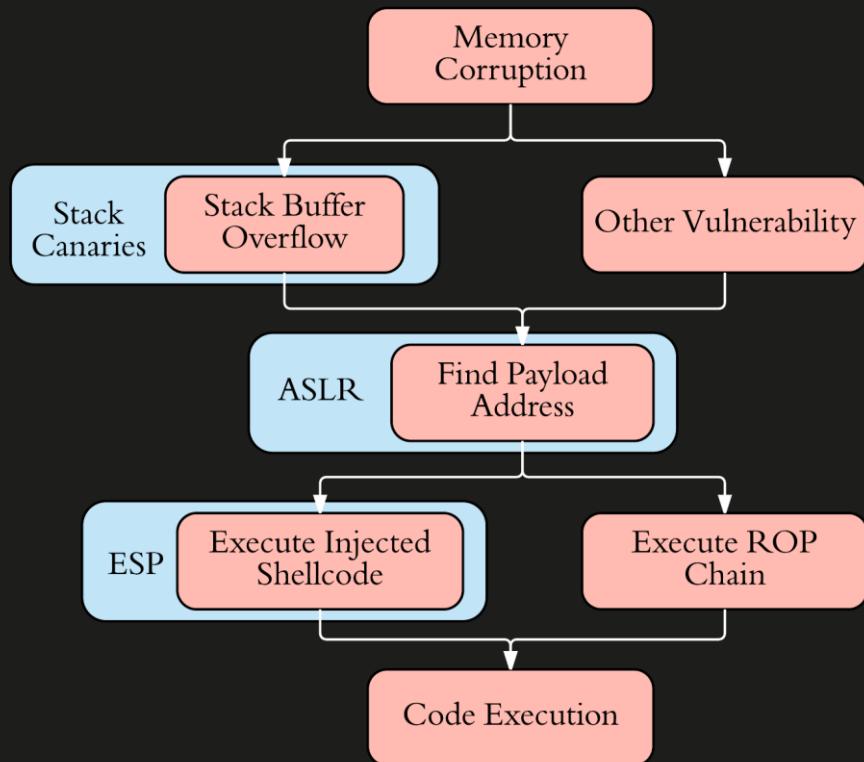
¹ XM6180 on iPhone4 has DEP, see '*All Your Baseband Are Belong To Us*'

² Newer Qualcomm chips run BLAST RTOS w. stack cookies / DEP / kernel-user separation, see '*Baseband Exploitation in 2013*'

Quantitative Analysis



Minimum Mitigation Baseline



- **ESP / DEP / NX / W^X**
Non-exec. data memory
- **ASLR**
Address Space Layout Randomization
- **Stack Canaries / SSP**
Stack buffer overflow protection



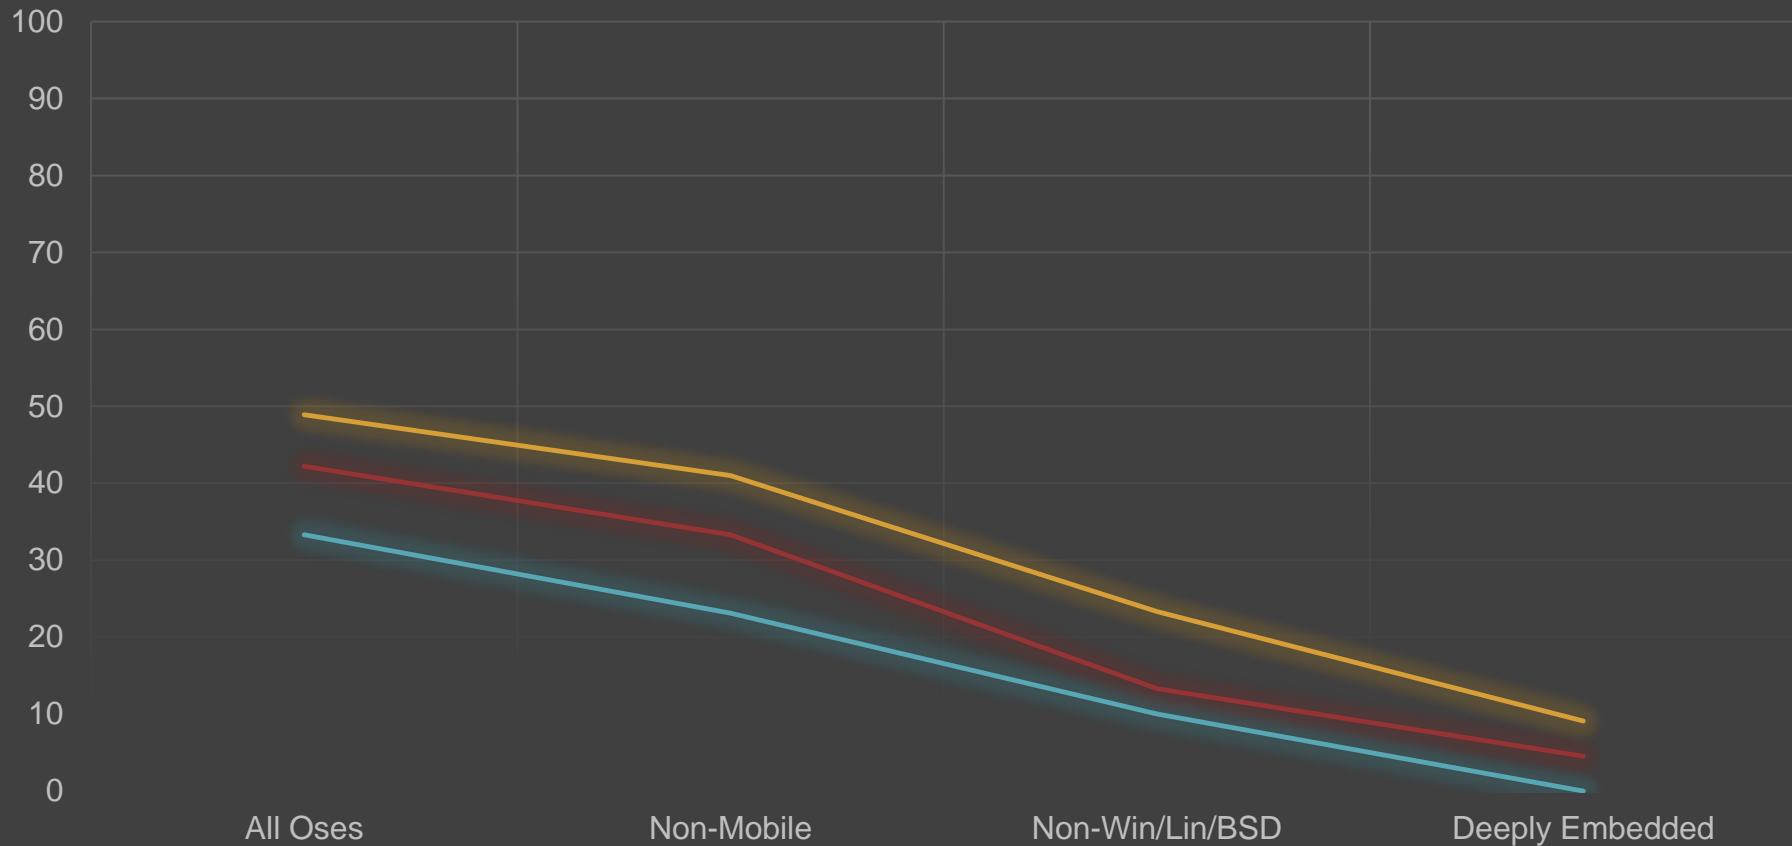
Operating System Selection

- Selected 45 Popular Embedded Oses
- High-end, Low-end, Linux/Windows/BSD-based, proprietary, etc.
- Evaluated Support For Mitigation Baseline

Optimistic Assesment

Mitigation Support

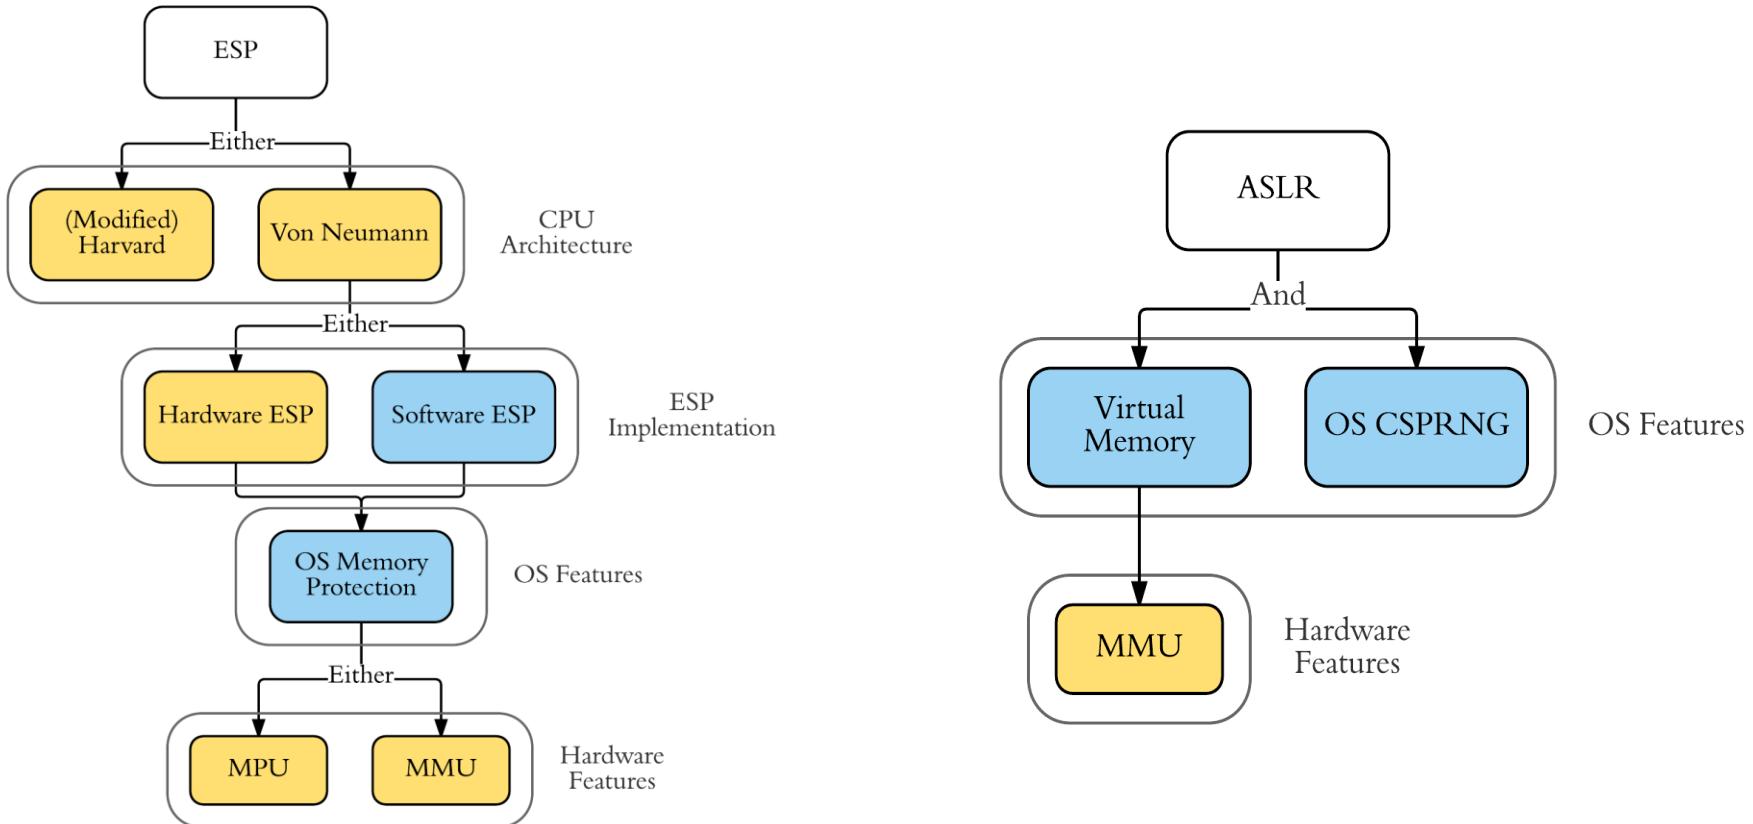
— ESP — ASLR — Stack Canaries



What's Going On Here?

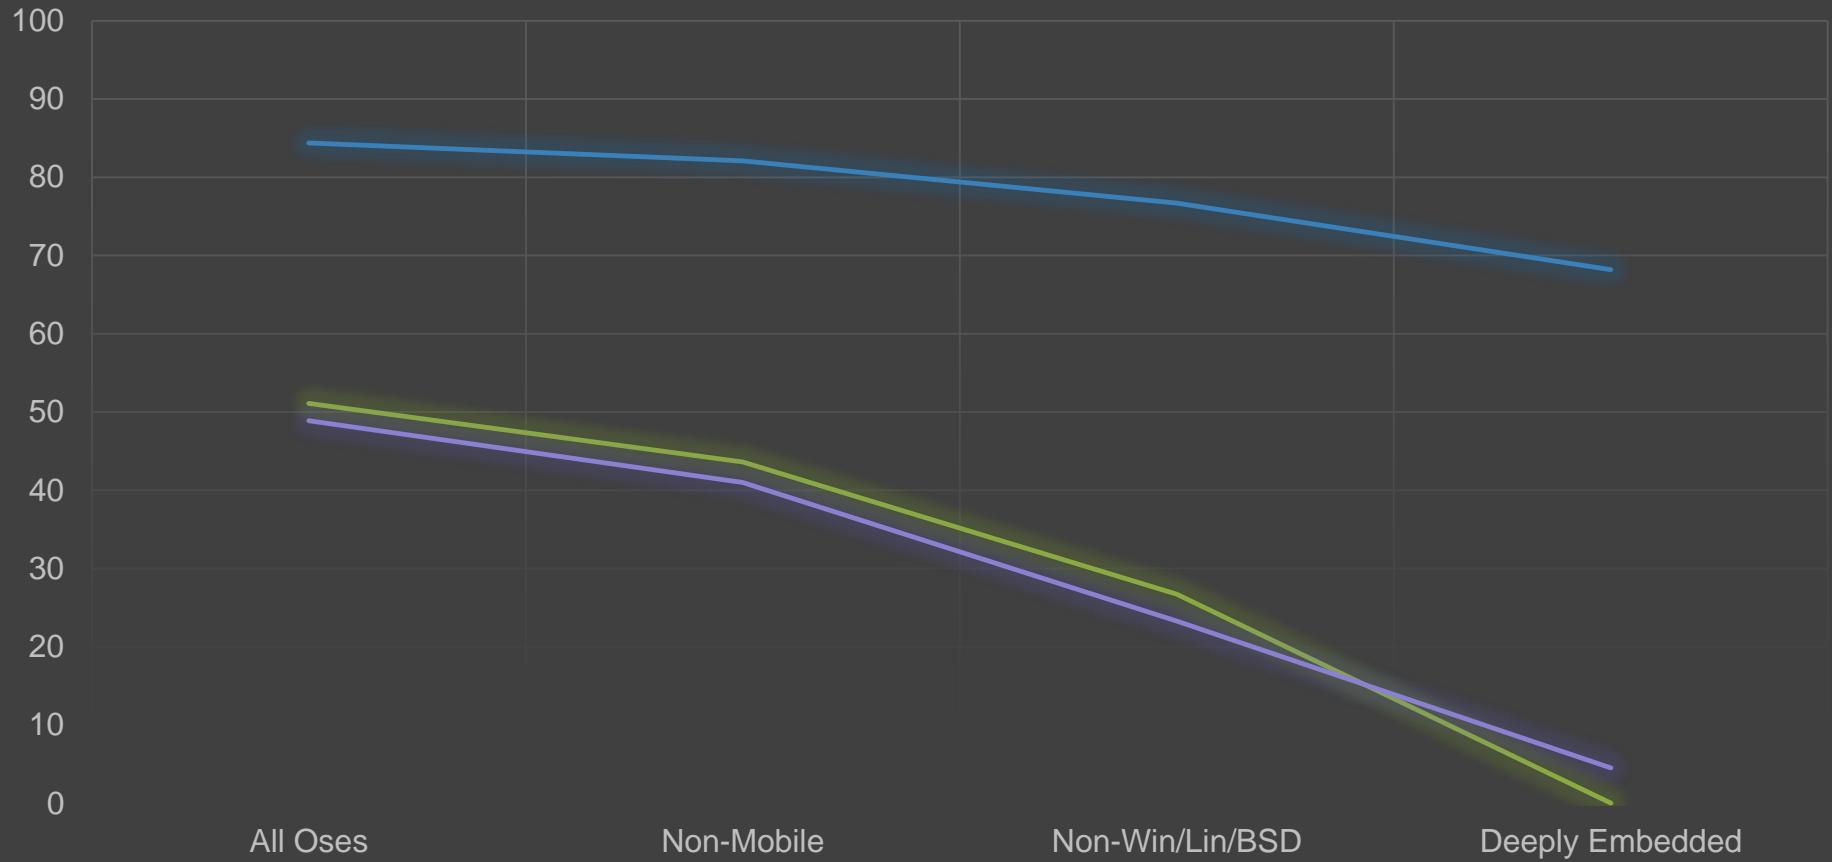


Hardware & Software Dependencies



Software Dependency Support

Memory Protection Virtual Memory OS CSPRNG



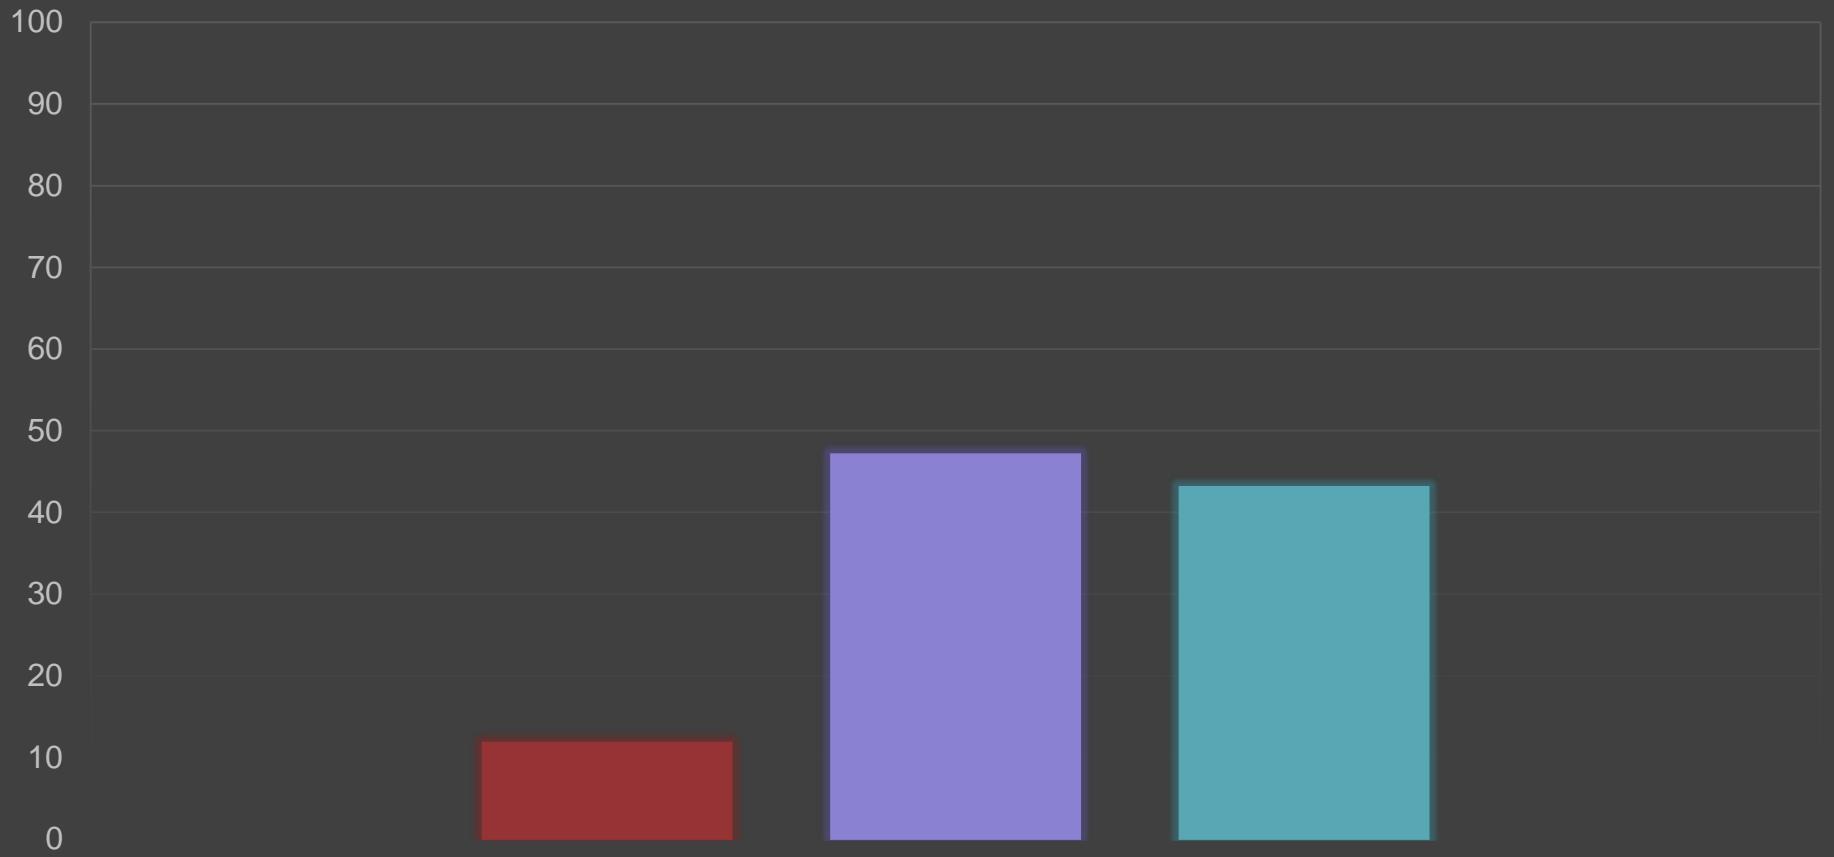


Hardware Selection

- Selected 78 Popular Embedded 'Core Families'
- Evaluated for Hardware Dependency Support

Hardware Dependency Support (VNA)

■ MPU ■ MMU ■ Hardware ESP



A photograph of a person from behind, wearing a green hooded jacket and a dark backpack, standing on a snowy mountain ridge. They are looking towards a range of snow-capped peaks under a blue sky with scattered white clouds.

Embedded Mitigation Challenges

Resource & System Constraints



- **Limited Resources**
Overhead influences adoption likelihood

- **Availability Requirements**
Guarantee high system uptime

- **Real-Time Conflicts**
Issues with eg. virtual memory caching, page faults, etc.



Advanced Processor Features

- Modern mitigations increasingly use advanced processor features eg. SGX, MPX, CET, LBR, PMU
- Only for modern high-end CPUs
- Embedded HW diversity -> OS cannot assume support

Embedded Random Number Generators

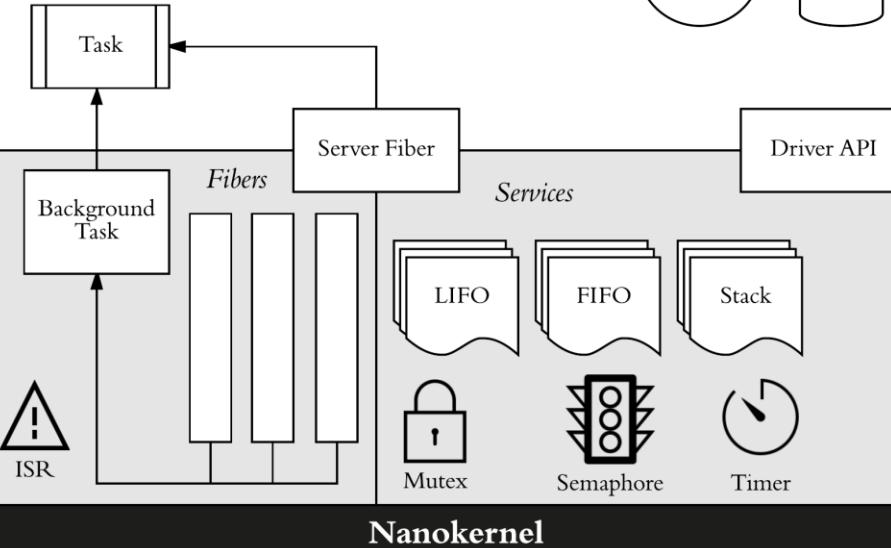
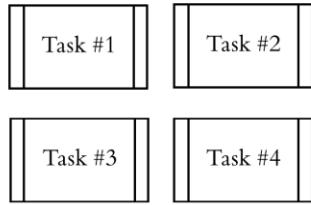


- Secure Randomness is OS Service!
- Hardware Diversity -> Little Entropy Source Omnipresence
- Low Entropy Environment
 - Little / predictable activity
 - "*Boot-time Entropy Hole*"
 - Can't wait for entropy too long..
- Result: not-so-random NGs, bad workarounds*

A black and white photograph showing the aftermath of a disaster, likely an earthquake or bombing. In the foreground, there is a large pile of twisted metal and rubble. Behind it, a city skyline is visible, with many multi-story buildings completely collapsed, their frames twisted and broken. The sky is overcast and hazy.

RTOS Mitigation Examples

Microkernel



- **Library based RTOS**
Based on Wind River Rocket
- **OS Linux Foundation Project**
Aimed at resource-constrained IoT
- **Young (2016) but promising**
Input from major chipmakers

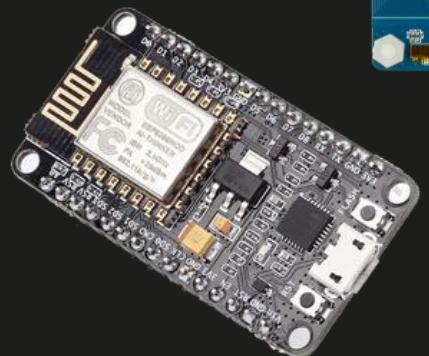
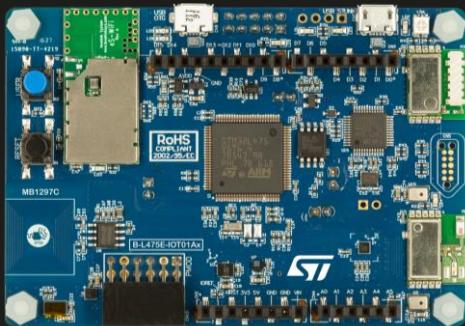
Explicit focus on security



Zephyr Stack Canaries

```
349 FUNC_NORETURN void _Cstart(void)
350 {
351 #ifdef CONFIG_ARCH_HAS_CUSTOM_SWAP_TO_MAIN
352     struct k_thread *dummy_thread = NULL;
353 #else
354     struct k_thread dummy_thread_memory;
355     struct k_thread *dummy_thread = &dummy_thread_memory;
356 #endif
357
358     /*
359     * Initialize kernel data structures. This step includes
360     * initializing the interrupt subsystem, which must be performed
361     * before the hardware initialization phase.
362     */
363
364     prepare_multithreading(dummy_thread);
365
366     /* perform basic hardware initialization */
367     _sys_device_do_config_level(_SYS_INIT_LEVEL_PRE_KERNEL_1);
368     _sys_device_do_config_level(_SYS_INIT_LEVEL_PRE_KERNEL_2);
369
370     /* initialize stack canaries */
371 #ifdef CONFIG_STACK_CANARIES
372     __stack_chk_guard = (void *)sys_rand32_get();
373 #endif
374
375     /* display boot banner */
376
377     switch_to_main_thread();
378
379     /*
380     * Compiler can't tell that the above routines won't return and issues
381     * a warning unless we explicitly tell it that control never gets this
382     * far.
383     */
384
385     CODE_UNREACHABLE;
386 }
```

- Based on Clang/GCC SSP
- Canary Failure -> System Error
- One master canary for entire address space
 - Generated once at system boot
 - No refreshing
- Generated using RNG API
 - Implementation depends on chosen *random* driver



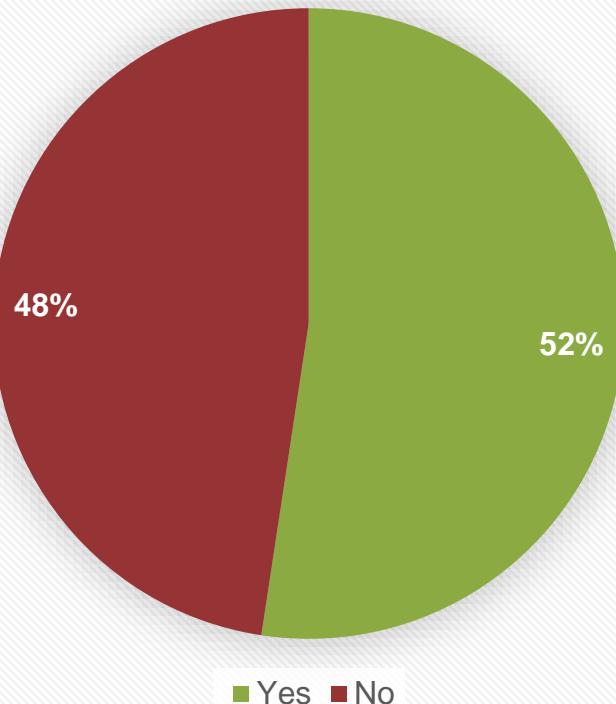
Zephyr Random API

RANDOM_HAS_DRIVER (TRNG)

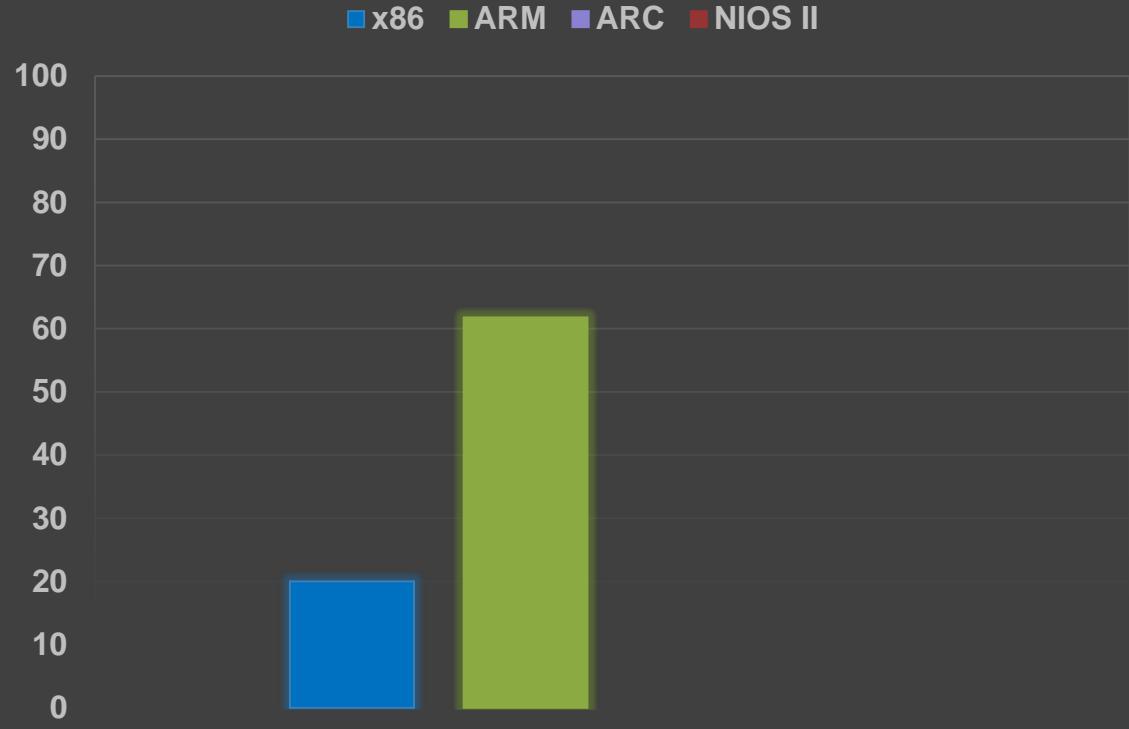
- **RANDOM_MCUX_RNGA**
NXP Kinetis K64F
- **RANDOM_MCUX_TRNG**
NXP Kinetis KW40Z & KW41Z
- **RANDOM_STM32_RNG**
STM32 Boards
- **RANDOM_ESP32_RNG**
ESP32 Boards
(requires Wi-Fi & Bluetooth enabled)

TRNGs Among Zephyr 1.8 Supported Boards

TRNG Support



TRNG Support Per Architecture





Zephyr Random API **TEST_RANDOM_DRIVER (PRNG)**

- **X86_TSC_RANDOM_GENERATOR**
Uses x86 RDTSC
 - Canary 1st value drawn from PRNG
little difference between bootruns
 - Infoleaks everywhere
- **TIMER_RANDOM_GENERATOR**
Suffers from similar issues
- **Had contact with Zephyr team,
plans to integrate OS CSPRNG**
Most likely NIST SP800-90A based



Fuchsia OS

- **Under-development Google RTOS**
Pushed to github w/o official statement in August 2016
- **Seems intended for IoT to smartphones**
- **Based on Magenta microkernel**
Derived from LK (Little Kernel)
Virtual Memory / MMU support
Kernel- / Userspace Separation

```

} regions[] = {
{
    .name = "kernel_code",
    .base = (vaddr_t)&__code_start,
    .size = ROUNDUP((size_t)&__code_end - (size_t)&__code_start, PAGE_SIZE),
    .arch_mmu_flags = ARCH_MMU_FLAG_PERM_READ | ARCH_MMU_FLAG_PERM_EXECUTE,
},
{
    .name = "kernel_rodata",
    .base = (vaddr_t)&__rodata_start,
    .size = ROUNDUP((size_t)&__rodata_end - (size_t)&__rodata_start, PAGE_SIZE),
    .arch_mmu_flags = ARCH_MMU_FLAG_PERM_READ,
},
{
    .name = "kernel_data",
    .base = (vaddr_t)&__data_start,
    .size = ROUNDUP((size_t)&__data_end - (size_t)&__data_start, PAGE_SIZE),
    .arch_mmu_flags = ARCH_MMU_FLAG_PERM_READ | ARCH_MMU_FLAG_PERM_WRITE,
},
590 void VmAspace::InitializeAslr() {
591     aslr_enabled_ = is_user() && !cmdline_get_bool("aslr.disable", false);
592
593     crypto::GlobalPRNG::GetInstance()->Draw(aslr_seed_, sizeof(aslr_seed_));
594     aslr_prng_.AddEntropy(aslr_seed_, sizeof(aslr_seed_));
595 }
596
597 uintptr_t choose_stack_guard(void) {
598     uintptr_t guard;
599     if (hw_rng_get_entropy(&guard, sizeof(guard), true) != sizeof(guard)) {
600         // We can't get a random value, so use a randomish value.
601         guard = 0xdeadbeef00ff00ffUL ^ (uintptr_t)&guard;
602     }
603     return guard;
}

```

Fuchsia OS Exploit Mitigations

- ESP / DEP

x86 MMU-based DEP policies

- ASLR

All executables PIE by default

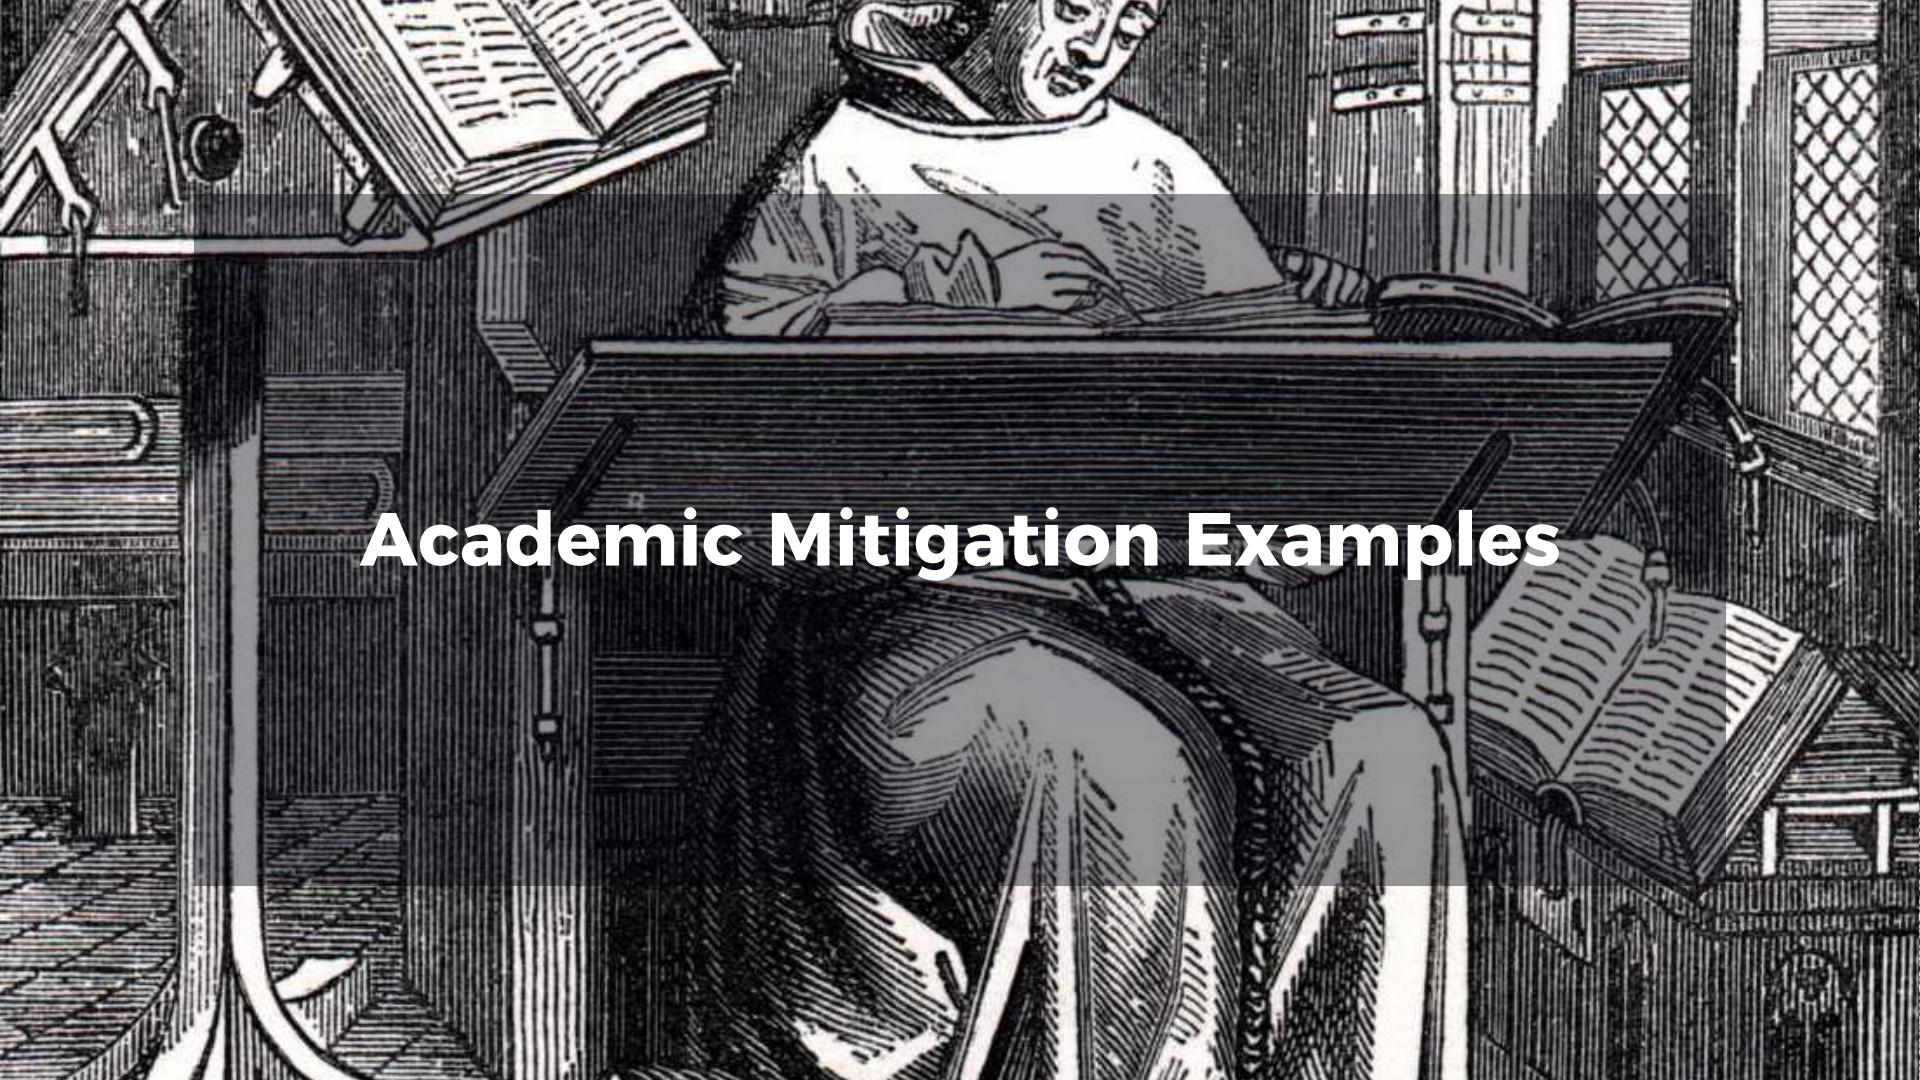
Randomization draws on OS CSPRNG

- Stack Canaries (GCC SSP Model)

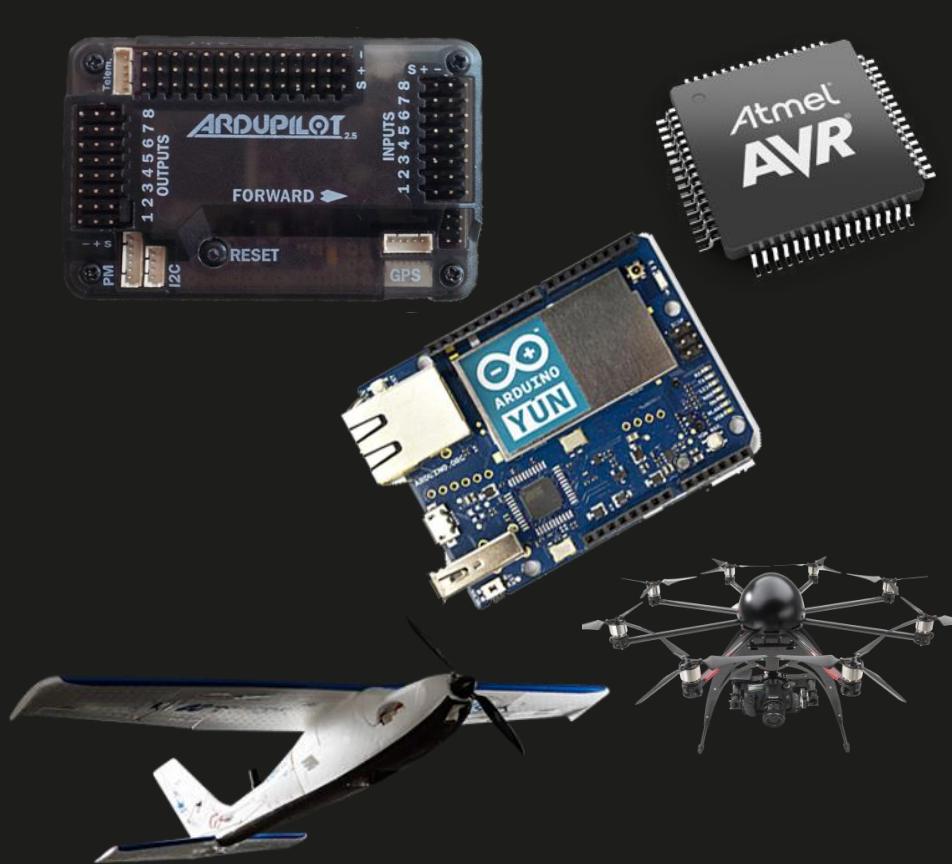
Single early-boot global canary

Canary generated using HW RNG API
Static value fallback

Canary failure -> kernel panic



Academic Mitigation Examples



* MAVR: *Code Reuse Stealthy Attacks and Mitigation on Unmanned Aerial Vehicles* – Habibi et al. (2015)

** AVRAND: *A Software-Based Defense Against Code Reuse Attacks for AVR Embedded Devices* – Pastrana et al. (2016)

MAVR* & AVRAND**

- Boot-Time Code Diversification
- MAVR: Ardupilot UAV system
- AVRAND: AVR chip in Arduino Yun
- **MAVR requires hardware mod**
Additional master processor + flash chip
- **(Re)Randomize by re-flashing**
Every N restarts or after crash
- **Attacker can reduce device lifespan by forcing re-flash to exhaust program / erase cycles**

A close-up photograph of a person's hand holding a silver pen, writing in a white notebook with blue horizontal grid lines. In the background, another page of the notebook shows a list of items with blue checkboxes next to them, such as "Drive" and "M2".

Deeply Embedded Exploit Mitigation Criteria



Embedded Exploit Mitigation Criteria

- **Low Resource Pressure**
Max 5% Overhead*
- **Hardware Agnostic**
- **Availability Preservation**
- **'Real-Time Friendly'**
- **Easy RTOS Integration**

* Microsoft BlueHat Prize Rules (2012),
'SoK: Eternal War in Memory' – Szekeres et al. (2013)

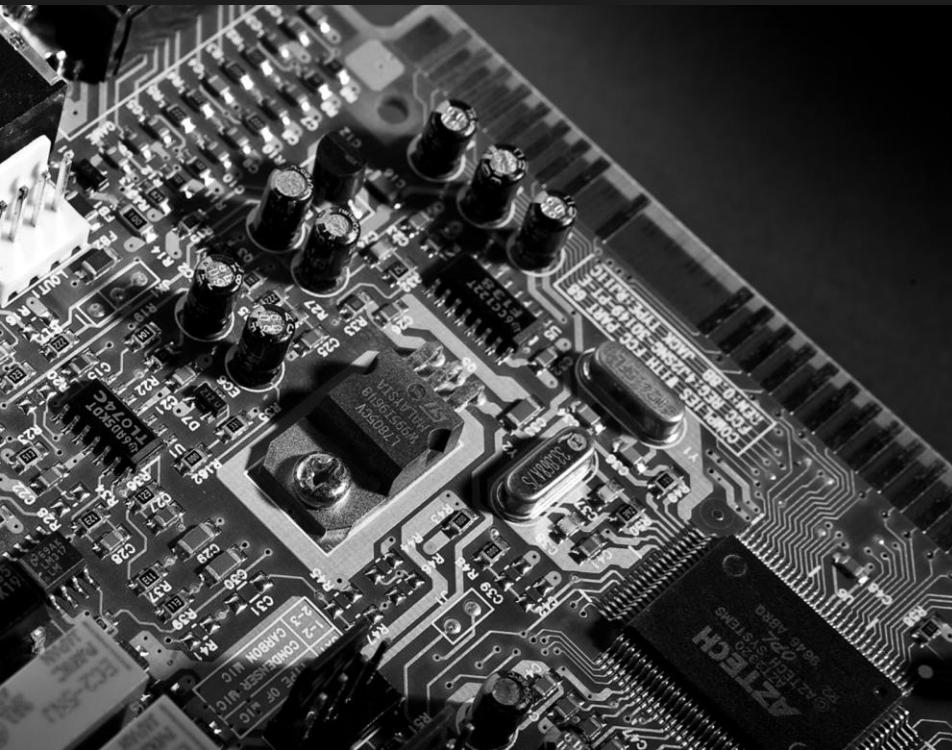


Embedded OS CSPRNG Criteria

- **Low Resource Pressure**
Use lightweight crypto
- **Entropy Gathering Limitations**
Rapidly available upon boot
No constant runtime polling
- **Non-Domain Specific Entropy Sources**
No reliance on sensor values,
accelerometers, etc.



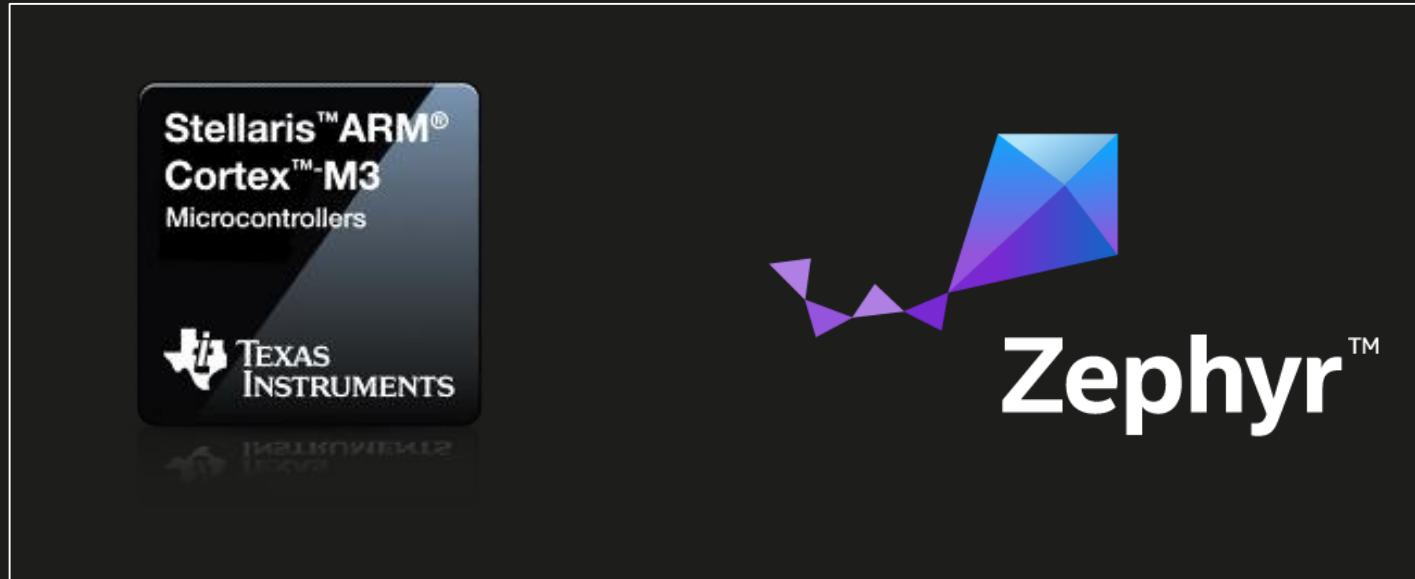
μArmor



µArmor

- **Deeply Embedded Mitigation Baseline**
µESP + µScramble + µSSP + µRNG
- **Target Systems**
Harvard or VNA CPU w. HW ESP
Single-address space (RT)OS or bare metal
- **Attacker Model**
Remote (eg. WiFi, Bluetooth, RF) control-flow hijacking attacks

Representative Platform



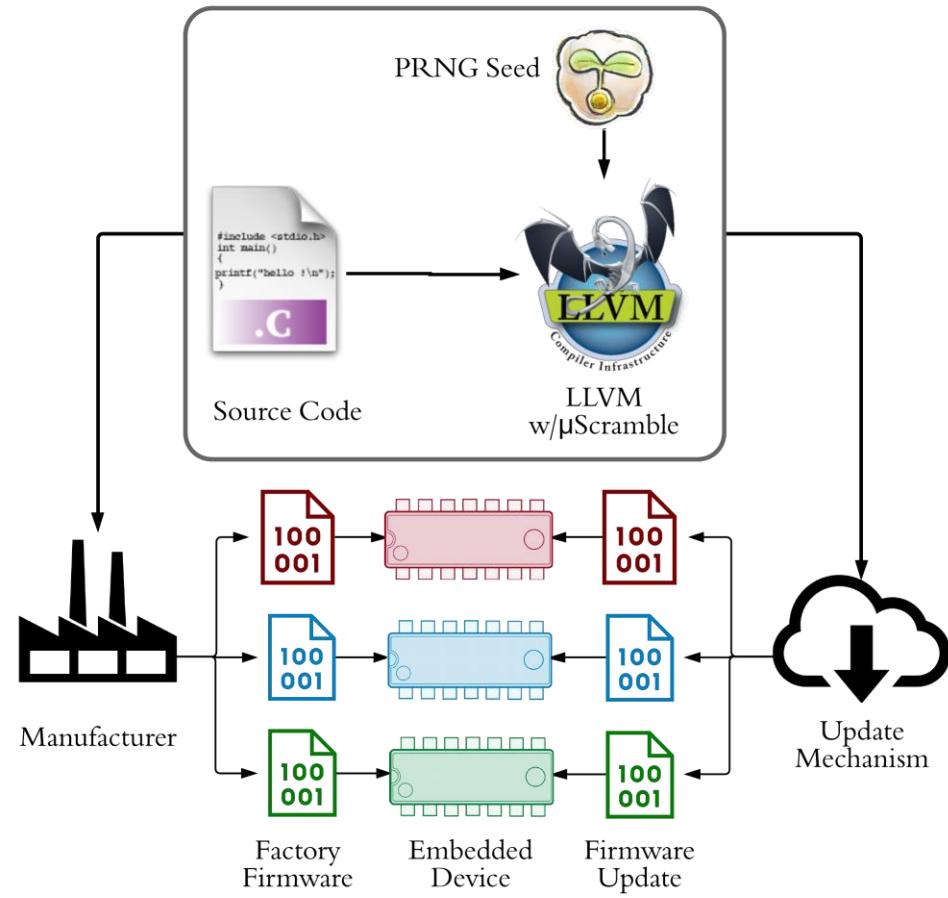
µESP

Memory	Permissions
Code (sensitive)	RO+XN
Code (other)	RO+X
Data	RW+XN
Peripherals	RW+XN
SCB Config	RO+XN
MPU Config	RO+XN

Table 14: µESP Memory Permission Policies

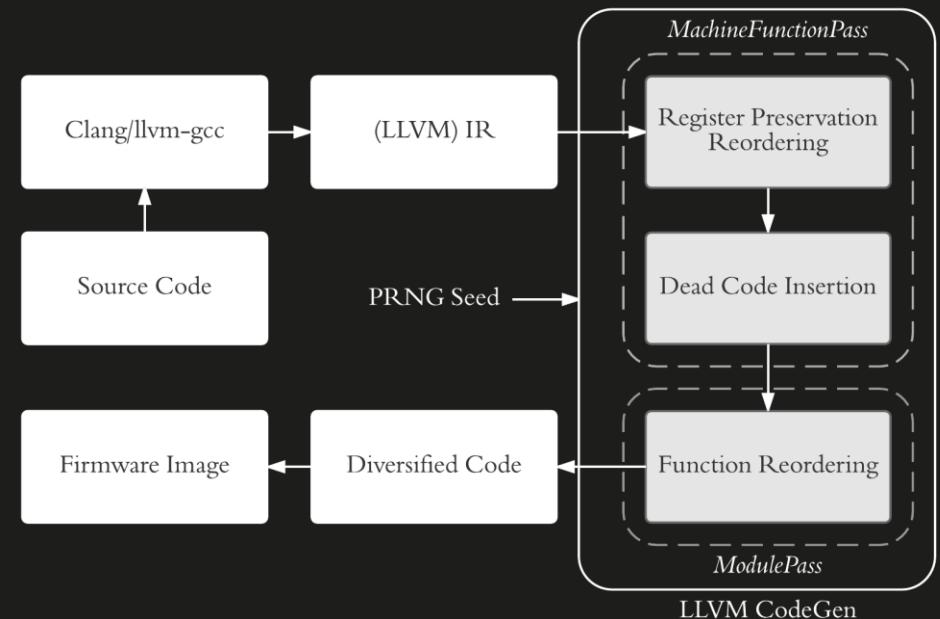
- **Use MPU to Enforce ESP**
- **Stack Grows Away From Data**
- **Limit Code-Reuse Surface**
Prevent ret2bootloader style attacks
- **'Lock Down' MPU & SCB**
Prevent MPU disabling

μScramble



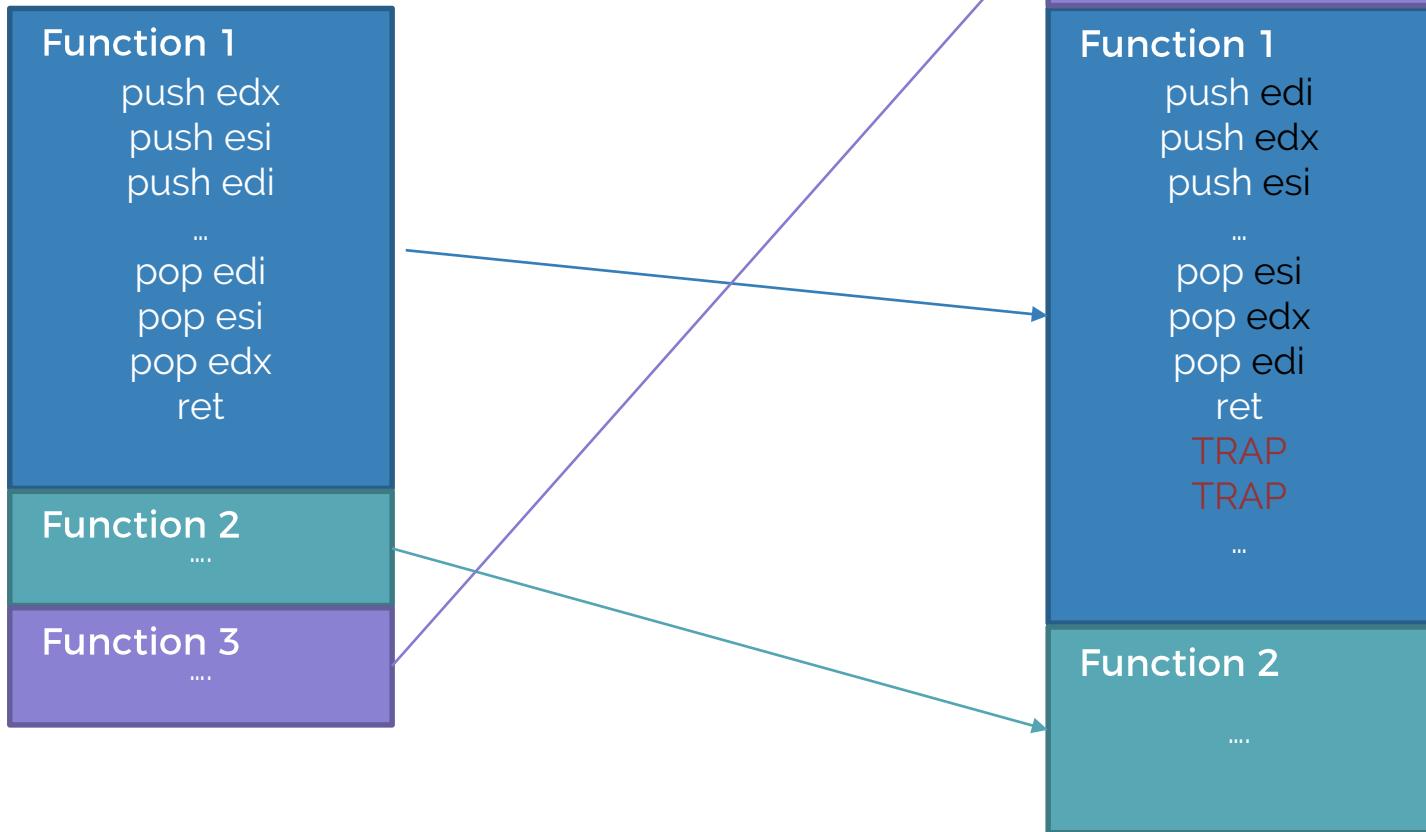
- Need Alternative for ASLR
- Software Diversification
 - When: Compilation, Boot, Runtime, Etc.
 - What: Code, Data
- Hybrid Compile/Update Solution
 - Integrated in firmware generation
 - Per-device diff firmware image
- Benefits
 - Limited Overhead
 - No virtual memory required
 - Works on source: platform independent

µScramble Code Diversification

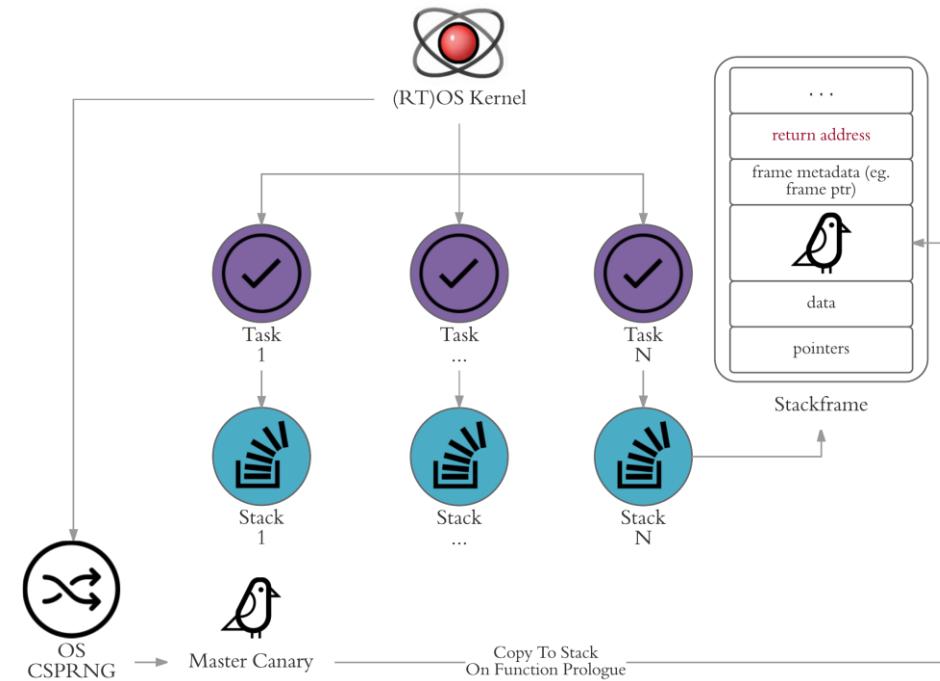


- **Code Transformations**
Function Reordering
Dead Code Insertion
Register Preservation Reordering
- **Fine-Grained Randomization**
Wrt. image base
Wrt. function base
Wrt. inter-function distance
Order of register preservation code
- **Decorrelation reduces infoleak & bruteforce impact**

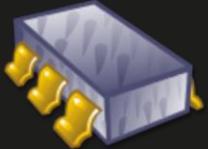
µScramble Transformations



µSSP

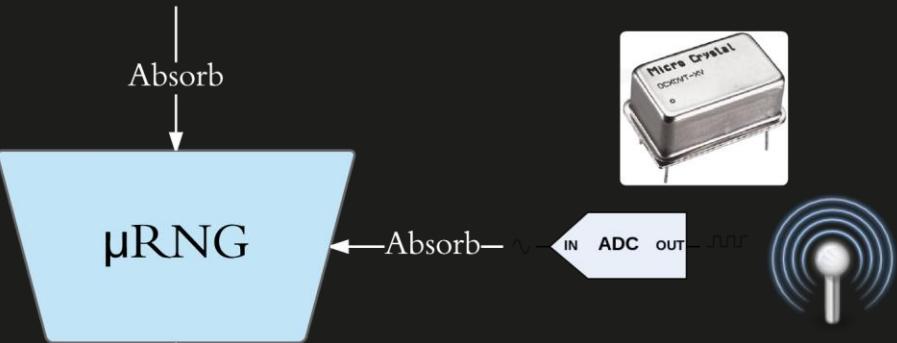


- **Based on Zephyr SSP**
- **Master canary drawn from OS CSPRNG**
- **Modular violation handler**
 - Passive
 - Fatal
- Thread Terminate / Restart
- System Terminate / Restart



Initial Entropy

Absorb



Pseudo-Random
Bits

μRNG

- **Keccak-based OS CSPRNG**
128-bit security

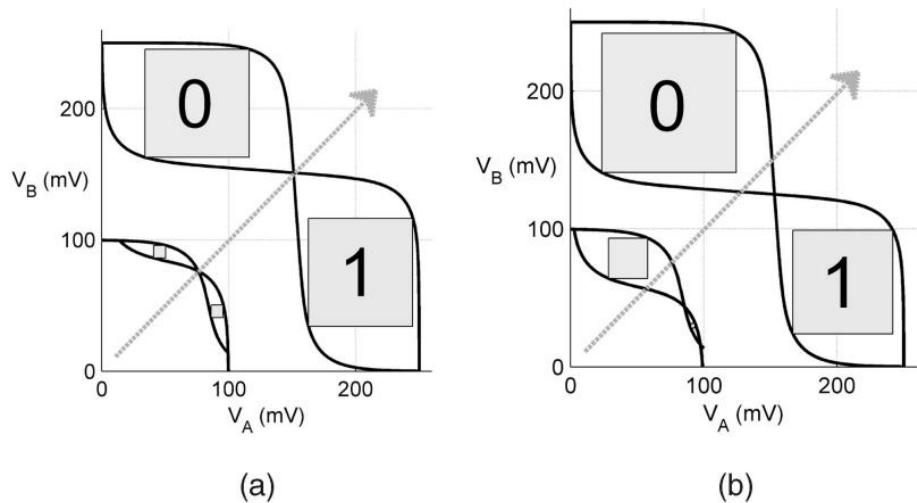
25 byte internal state

1 algorithm for entropy gather & RNG

Integrated as Zephyr random driver

- **Avoid Reseed Polling**
Reseed Control linked to PRNG output

μ RNG Initial Entropy Source



- **SRAM Startup Values (SUVs)**
Two stable states
Converge after initial instability

 - **Manufacturing Differences**
Biased cells -> PUFs
Random cells -> PRNGs

 - **Many Benefits**
Instantly available at very early boot
- Differs between devices & bootruns

SRAM omnipresent on MCUs



µArmor

Overhead & Security Evaluation

µArmor

Overhead Evaluation

RNG	ROM	RAM
µRNG ¹	709	81
TinyRNG [302] ²	11086	471
TinyKey [428] ³	5990	790

- **Metrics**

Code Size, Data Size, Memory Usage, Runtime Overhead

- **Evaluated Against**

IoT sample applications

TACLeBench suite

- **Overhead**

Code Size $\leq 5\%$

Others $\leq 1\%$

Set	Avg. GS ¹	Max. GS ¹
Sequential		
adpcm_dec	0.5	4
adpcm_enc	0.6	5
ammunition	3.6	147
anagram	0.6	4
audiobeam	4	79
cjpeg_transupp	32	96
cjpeg_wrbmp	1.4	3
dijkstra	1.3	5
epic	0	0
fref	0.5	3
gsm_dec	1.3	43

µScramble Security Evaluation

- **Coverage Analysis**

1000 firmware variants per benchmark

Harvest ROP gadgets

Max avg. gadget survival rate: 10.15%

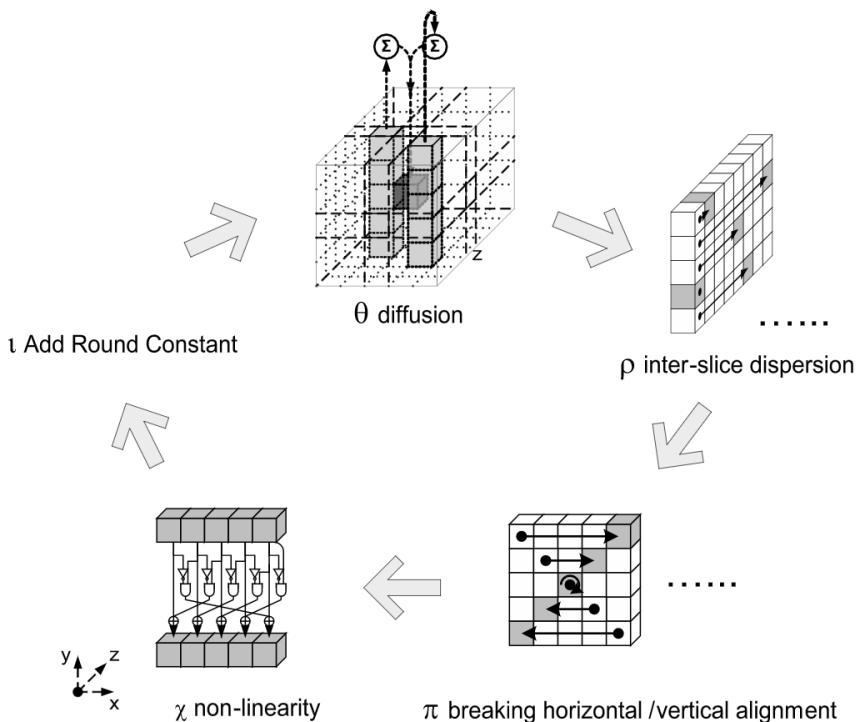
- **Complicating Attacker Factors**

Payload length

Firmware complexity

Fine-Grained Randomization

μ RNG Security Evaluation



- Keccak is well-analyzed
- **Passive State Recovery Attacks**
Reseed at least every 1GB of output
well within 32GB bound*
- **No 'Boot Time Entropy Hole'**
All initial entropy instantly available

µRNG

Security Evaluation

Table 1. Min-entropy results for STM32F100R8 SRAMs at different temperatures. Min-entropy denoted as percentage of total available SRAM.

Temp. [°C]	Microcontroller ID									
	1	2	3	4	5	6	7	8	9	10
-30	5.3%	5.3%	5.4%	5.5%	5.4%	5.3%	5.4%	5.2%	5.3%	5.8%
+25	6.6%	6.6%	6.7%	6.8%	6.7%	6.5%	6.8%	6.5%	6.7%	6.7%
+90	6.3%	6.5%	6.5%	6.6%	6.2%	6.5%	6.5%	6.2%	6.5%	6.5%

- **SRAM Suitability Depends on Chip**
Need Entropy $\geq 2^*$ security strength

Avg. Entropy: 5% of SRAM size*
Worst-case: 1.2% (100 bits) for PIC16F

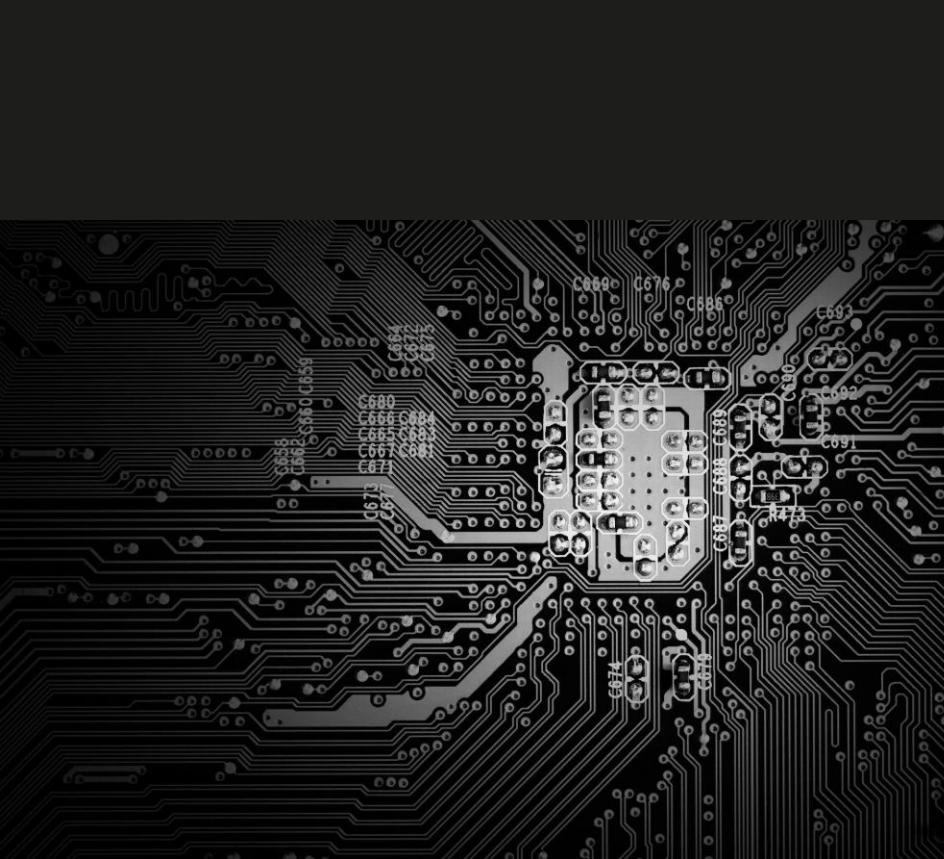
- **Reality**
SRAM slightly contaminated
Harvest in v. early boot & limit stack use

- **Memory Retention**
Minimal at normal ambient temps
Clear loss in cold (-20C) for some chips

Full SRAM reset between power cycles

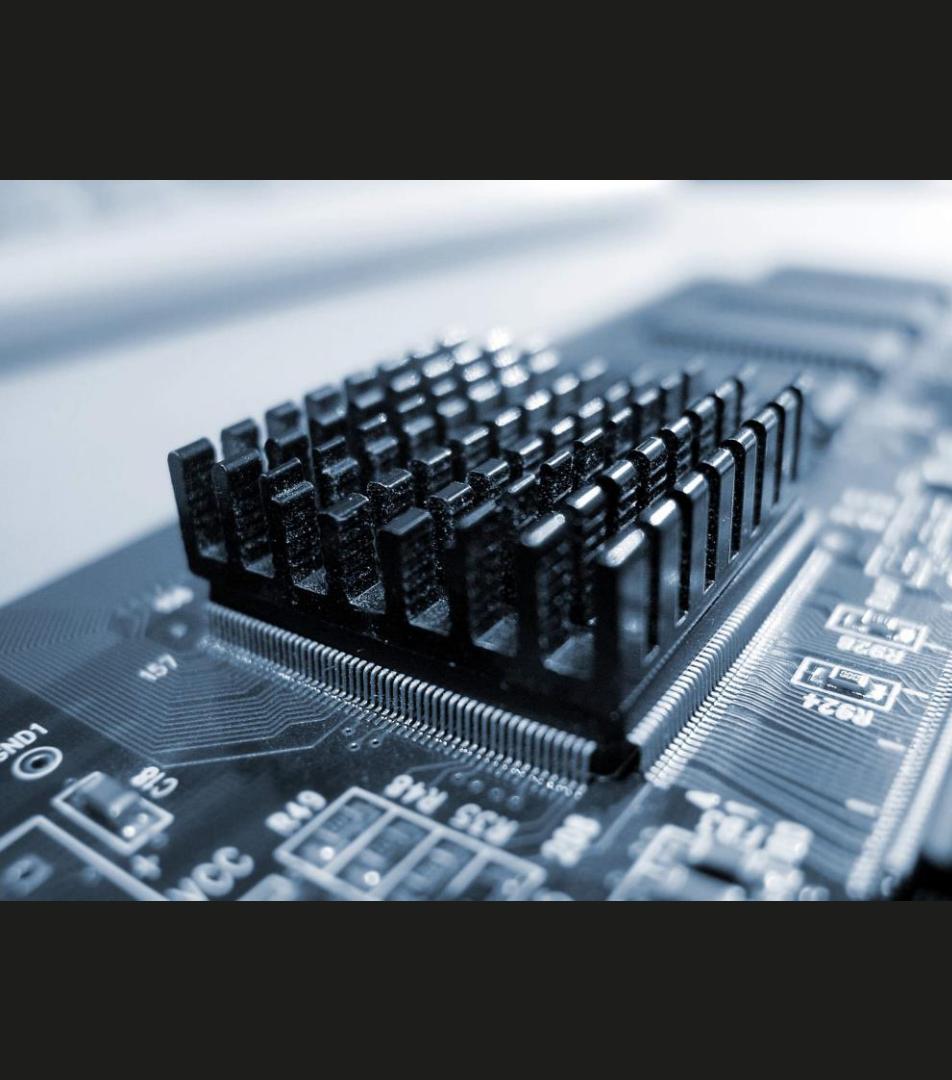
* Secure PRNG Seeding on Commercial Off-the-Shelf Microcontrollers -
Van Herrewege et al. (2013)

Limitations, Conclusions & Future Work



uArmor Limitations

- **Platform**
Harvard CPU or VNA with hardware ESP
- **Attacker Model**
Remote Control-Flow Hijacking Only
- **Diversification Relocation Frequency**
Compile- / Update-Time Only
- **Diversification Coverage**
Code Memory only

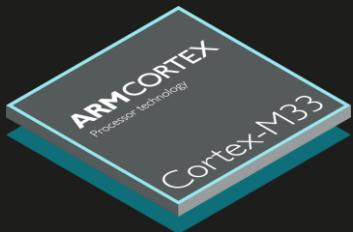


Conclusions

- **Embedded memory Corruption issues warrant serious attention**
Smallest of devices are connected
Unsafe languages are here to stay
- **Current state of embedded mitigations is terrible**
Esp. in deeply embedded / RTOS
Not easily fixed either
- **µArmor is intended as small step forward**



Tock



Future Work

- **Long Term**

Embedded Safe Language Adoption
(Rust Embedded, Tock OS, etc.)

Secure & Scalable Patching

- **Short Term**

Mitigation Baseline Adoption in RTOS

Hardware Dependency Support

Advanced Embedded Mitigations
(CFI, XoM, etc.)

Questions?

 @s4mvartaka
 j.wetzels@midnightbluelabs.com

