

Hardware implementation of Keyak₂

Jos Wetzels & Wouter Bokslag

Content

- Authenticated Encryption
- CAESAR Competition
- Keyak₂
- Scope
- GMU API
- Design Decisions
- Test Benches
- Performance Figures

Authenticated Encryption

- Encryption provides confidentiality but not (per se) *integrity* or *authenticity*
- Problem when attacker can modify ciphertext without detection (eg. padding oracle attacks, bitflipping, etc.)
- We want authentication of our data:
 - AE: Authenticated Encryption
 - AEAD: Authenticated Encryption with Associated Data (also include plaintext metadata in check)
- Traditional solution: Message Authentication Codes (MACs)
- Implementation often goes wrong (see POODLE attack on TLS)
- Desire for more dedicated AEAD schemes outside of AES in GCM/CCM/OCB mode

CAESAR Competition

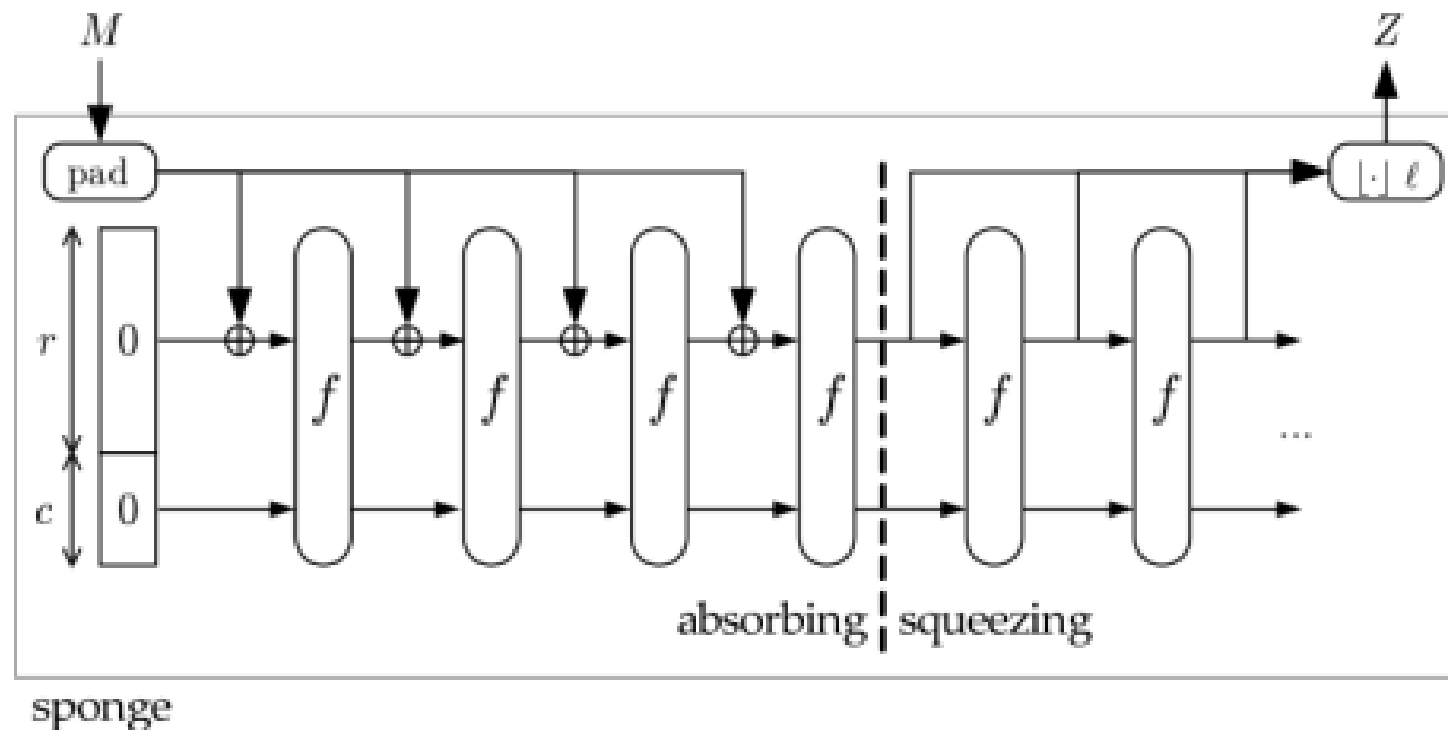
- Competition for Authenticated Encryption, selects portfolio of schemes
- Primary candidate goals:
 - Offer advantages over AES-GCM
 - Suitable for widespread adoption
- Candidates specify *family* of authenticated ciphers
- Five inputs, one output:
 - Input Data (eg. plaintext or ciphertext)
 - Associated Data (AD), also called metadata (eg. packet headers)
 - Key
 - (*optional*) Public Message Number
 - (*optional*) Secret Message Number
- Output Data (eg. ciphertext or decrypted plaintext)

Keyak₂ - Overview

- Parameterized permutation-based authenticated encryption scheme
- 2nd round CAESAR candidate
- 5 named instances aimed at spectrum of platforms
- Offers support for *session-based* authentication (full sequence of messages is authenticated)
- In-place encryption & decryption
- Stream-compatible (does not require prior knowledge of input length)

Keyak₂ - Sponges

- Mode of operation is 'sponge-based' construction



- Data is first 'absorbed' into sponge state, later 'squeezed' out
- Primitive suitable to make hashes, MACs, stream ciphers, etc.
 - eg. 'squeeze' keystream from it, absorb data in it for later authentication

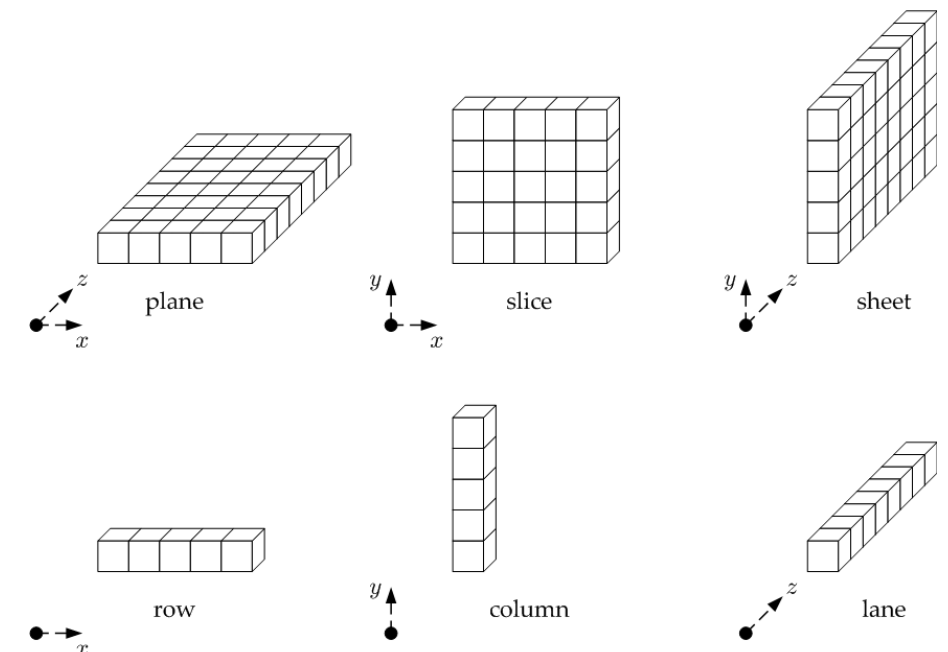
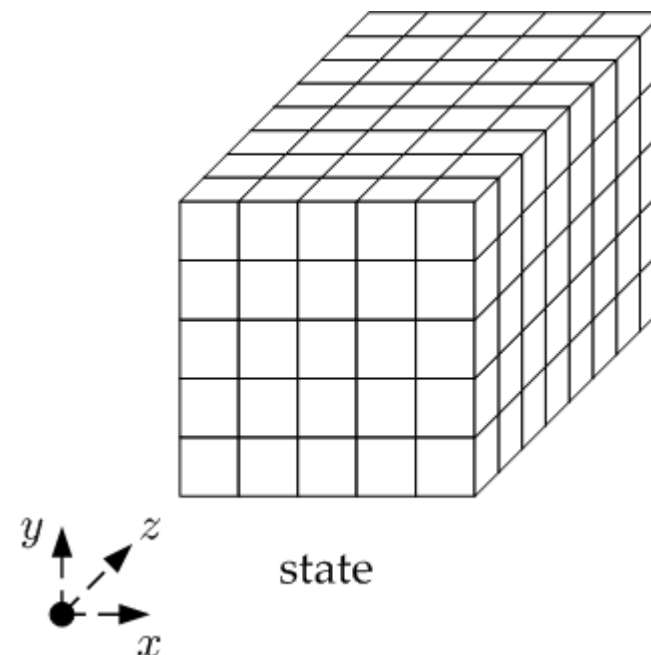
Keyak₂ – Motorist Mode

- New mode of operation (“*Motorist*”) compared to Keyak₁
- Consists of 3 components:
 - Piston: Has internal state, performs ‘sponge’ operations on it
 - *Crypt*: Encryption/Decryption, squeeze ‘keystream’ from state, absorb input into state
 - *Inject*: Absorb metadata into state
 - *Spark*: Perform permutation (core of scheme security)
 - *GetTag*: Extract authentication tag from state
 - Engine: Controls array of π parallel pistons, effectively wrapper around piston functionality
 - Motorist: Top-level component, performs initialization and ‘wrapping’
 - *StartEngine*: Initializes state using SUV (Secret and Unique Value) which is derived from key + nonce
 - *Wrap*: Performs AEAD call (encrypt or decrypt + authenticate data + metadata)

Keyak₂ - Permutation

- Uses Keccak_p (derivative of Keccak_f) as underlying permutation
- Keccak_f also underlies SHA-3 winner Keccak
- Difference Keccak_p and Keccak_f: starting round index
- State represented as 5x5 array of n-bit '*lanes*'

```
Keccak-f[b](S):  
  for  $i \in \{0, \dots, (n_r - 1)\}$ :  
     $S = \text{Round}[b](S, RC_i)$   
  Return S
```



Keyak₂ - Permutation

- Round function 'shuffles' state according to several criteria

```

Round[b](A, RC):
  //θ-step for bit diffusion
  for  $x \in \{0, \dots, 4\}$ :
     $C[x] = A[x,0] \text{ xor } A[x,1] \text{ xor } A[x,2] \text{ xor } A[x,3] \text{ xor } A[x,4]$ 

  for  $x \in \{0, \dots, 4\}$ :
     $D[x] = C[x-1] \text{ xor } \text{rot}(C[x+1], 1)$ 

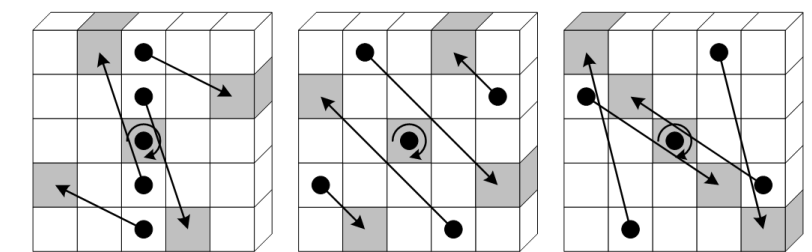
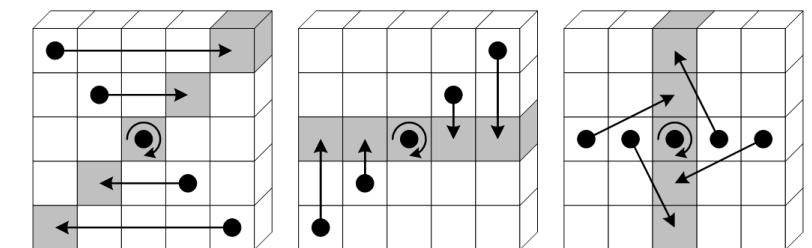
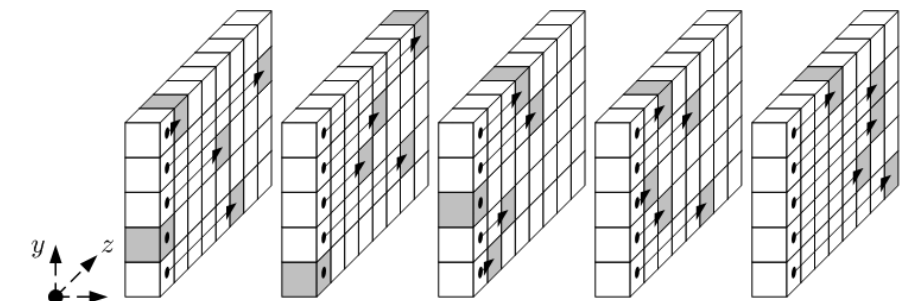
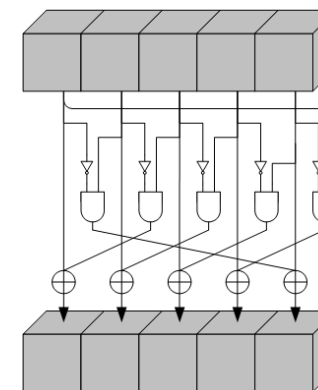
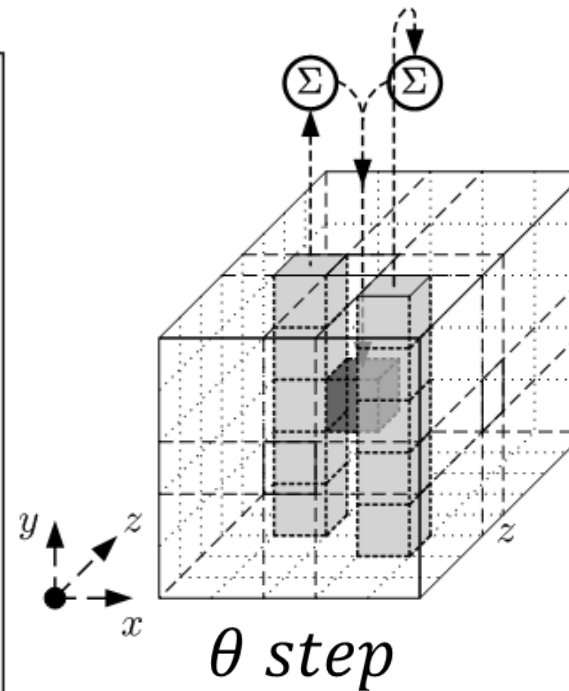
  for  $(x, y) \in \{\{0, \dots, 4\} \times \{0, \dots, 4\}\}$ :
     $A[x,y] = A[x,y] \text{ xor } D[x]$ 

  //ρ-step for inter-slice diffusion and
  //π-step for disturbing x,y alignment through lane-transposition
  for  $(x, y) \in \{\{0, \dots, 4\} \times \{0, \dots, 4\}\}$ :
     $B[y, 2x + 3y] = \text{rot}(A[x,y], r[x,y])$ 

  //χ-step for non-linear mapping
  for  $(x, y) \in \{\{0, \dots, 4\} \times \{0, \dots, 4\}\}$ :
     $A[x,y] = B[x,y] \text{ xor } ((\text{not } B[x+1,y]) \text{ and } B[x+2,y])$ 

  //ι-step to break symmetry
   $A[0,0] = A[0,0] \text{ xor } RC$ 

  Return A
  
```



Keyak₂ – Named Instances

- 5 Named Instances
- Differ in state-width b and piston-parallelism π
- Suitable for different platforms

RIVER KEYAK $b = 800, \pi = 1$

LAKE KEYAK $b = 1600, \pi = 1$ (primary recommendation)

SEA KEYAK $b = 1600, \pi = 2$

OCEAN KEYAK $b = 1600, \pi = 4$

LUNAR KEYAK $b = 1600, \pi = 8$

Keyak₂ – Pure Python Implementation

```
def _KeccakFonLanes(self, lanes):
    R = 1
    for roundIndex in range(self.aStartRoundIndex, (self.aStartRoundIndex + self.aNrRounds)):
        # θ
        C = [lanes[x][0] ^ lanes[x][1] ^ lanes[x][2] ^ lanes[x][3] ^ lanes[x][4] for x in range(5)]
        D = [C[(x+4)%5] ^ self._ROL(C[(x+1)%5], 1) for x in range(5)]
        lanes = [[lanes[x][y]^D[x] for y in range(5)] for x in range(5)]

        # ρ and π
        (x, y) = (1, 0)
        current = lanes[x][y]
        for t in range(24):
            (x, y) = (y, (2*x+3*y)%5)
            (current, lanes[x][y]) = (lanes[x][y], self._ROL(current, (t+1)*(t+2)//2))

        # χ
        for y in range(5):
            T = [lanes[x][y] for x in range(5)]
            for x in range(5):
                lanes[x][y] = T[x] ^ ((~T[(x+1)%5]) & T[(x+2)%5])

        # ι
        lanes[0][0] ^= self.RC[roundIndex]
    return lanes
```

```
221 def StartEngine(self, SUV, tagFlag, T, unwrapFlag, forgetFlag):
222     if (self.phase != MotoristPhase.ready):
223         raise Exception("The phase must be ready to call Motorist.StartEngine().")
224
225     self.engine.InjectCollective(SUV, True)
226
227     if (forgetFlag):
228         self._MakeKnot()
229
230     res = self._HandleTag(tagFlag, T, unwrapFlag)
231
232     if (res):
233         self.phase = MotoristPhase.riding
234     return res
235
236 def Wrap(self, I, O, A, T, unwrapFlag, forgetFlag):
237     if (self.phase != MotoristPhase.riding):
238         raise Exception("The phase must be riding to call Motorist.Wrap().")
239
240     if(not(hasMore(I)) and not(hasMore(A))):
241         self.engine.Inject(A)
242
243     while (hasMore(I)):
244         self.engine.Crypt(I, O, unwrapFlag)
245         self.engine.Inject(A)
246
247     while (hasMore(A)):
248         self.engine.Inject(A)
249
250     if ((self.Pi > 1) or (forgetFlag)):
251         self._MakeKnot()
252
253     res = self._HandleTag(True, T, unwrapFlag)
254
255     if not(res):
256         O.erase()
257     return res
```

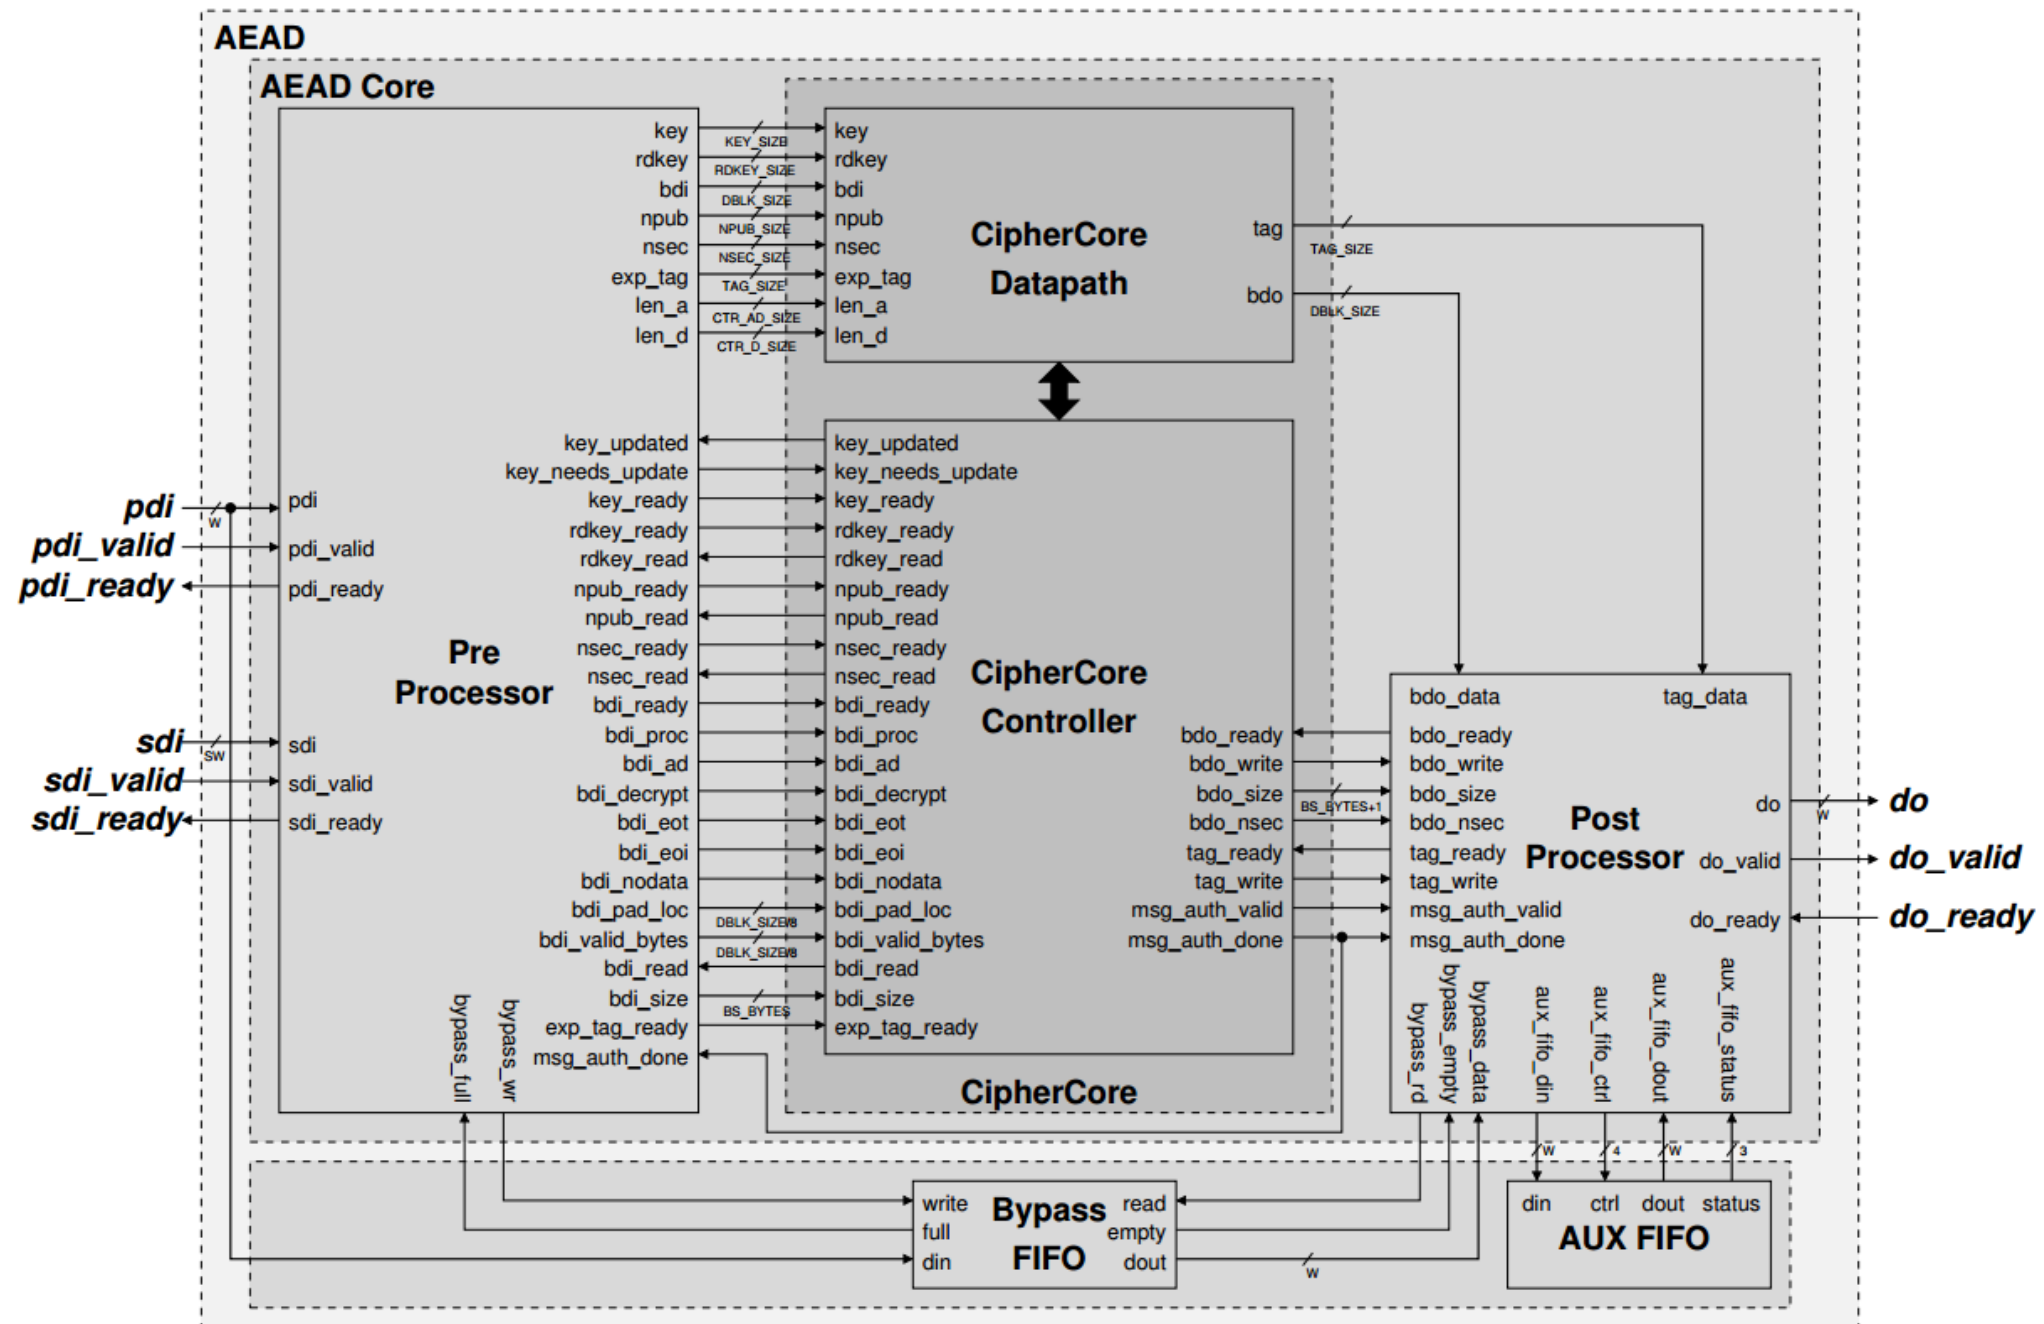
Internship Scope

- CAESAR 3rd round candidate process requires (reference) hardware implementation
- River & Lake Keyak most suitable (no piston-parallelism)
- River Keyak
 - Lane size: 32 bits
 - State width: 800 bits (5x5x32)
 - Squeeze rate: 68 bytes
 - Absorb rate: 96 bytes
- Lake Keyak
 - Lane size: 64 bits
 - State width: 1600 bits (5x5x64)
 - Squeeze rate: 168 bytes
 - Absorb rate: 192 bytes

GMU API

- CAESAR call for submissions only defined a software API
- 3rd round candidates require hardware implementation
- George Mason University (GMU) proposed hardware API
- Motivation:
 - Hardware API influences area & throughput/area ratio of implementation
 - Uniform hardware API facilitates candidate comparison
- Features:
 - Inputs of arbitrary number of bytes
 - Streaming-support: no need to know size of plain/cipher/AD in advance
 - Independent data & key inputs
 - Encryption/Decryption in same core, toggled with flag
 - External communication with AXI4 or FIFOs

GMU API



GMU API

- Input processing handled by PreProcessor
 - Key loading/activation
 - Data block loading, padding
 - Keeping track of data to process
 - Etc.
- Output processing handled by PostProcessor
 - Clearing output blocks not belonging to ciphertext or plaintext
 - Conversion of output blocks to data words
 - Holding decrypted plaintext in AUX FIFO buffer until full authentication
 - Generating errors upon authentication failure
 - Etc.
- Actual encryption/decryption handled by CipherCore
- Implementers only have to design CipherCore internals

Design Decisions

- Implementations concern named instance rather than family
- Hence can identify many simplifications, reductions, etc. in general algorithm (see design doc for exhaustive list and details)
- Several variables can be assume constant:
 - CAESAR requires static nonce, key & tag lengths. Use recommended values from Keyak specs:
 - Nonce: 464 or 1200 bits (River vs Lake)
 - Key: 128 bits
 - Tag: 128 bits
 - Sponge-related variables (squeeze & absorb rates, capacity, chaining value length)
 - Named instances have fixed state width & permutation round count

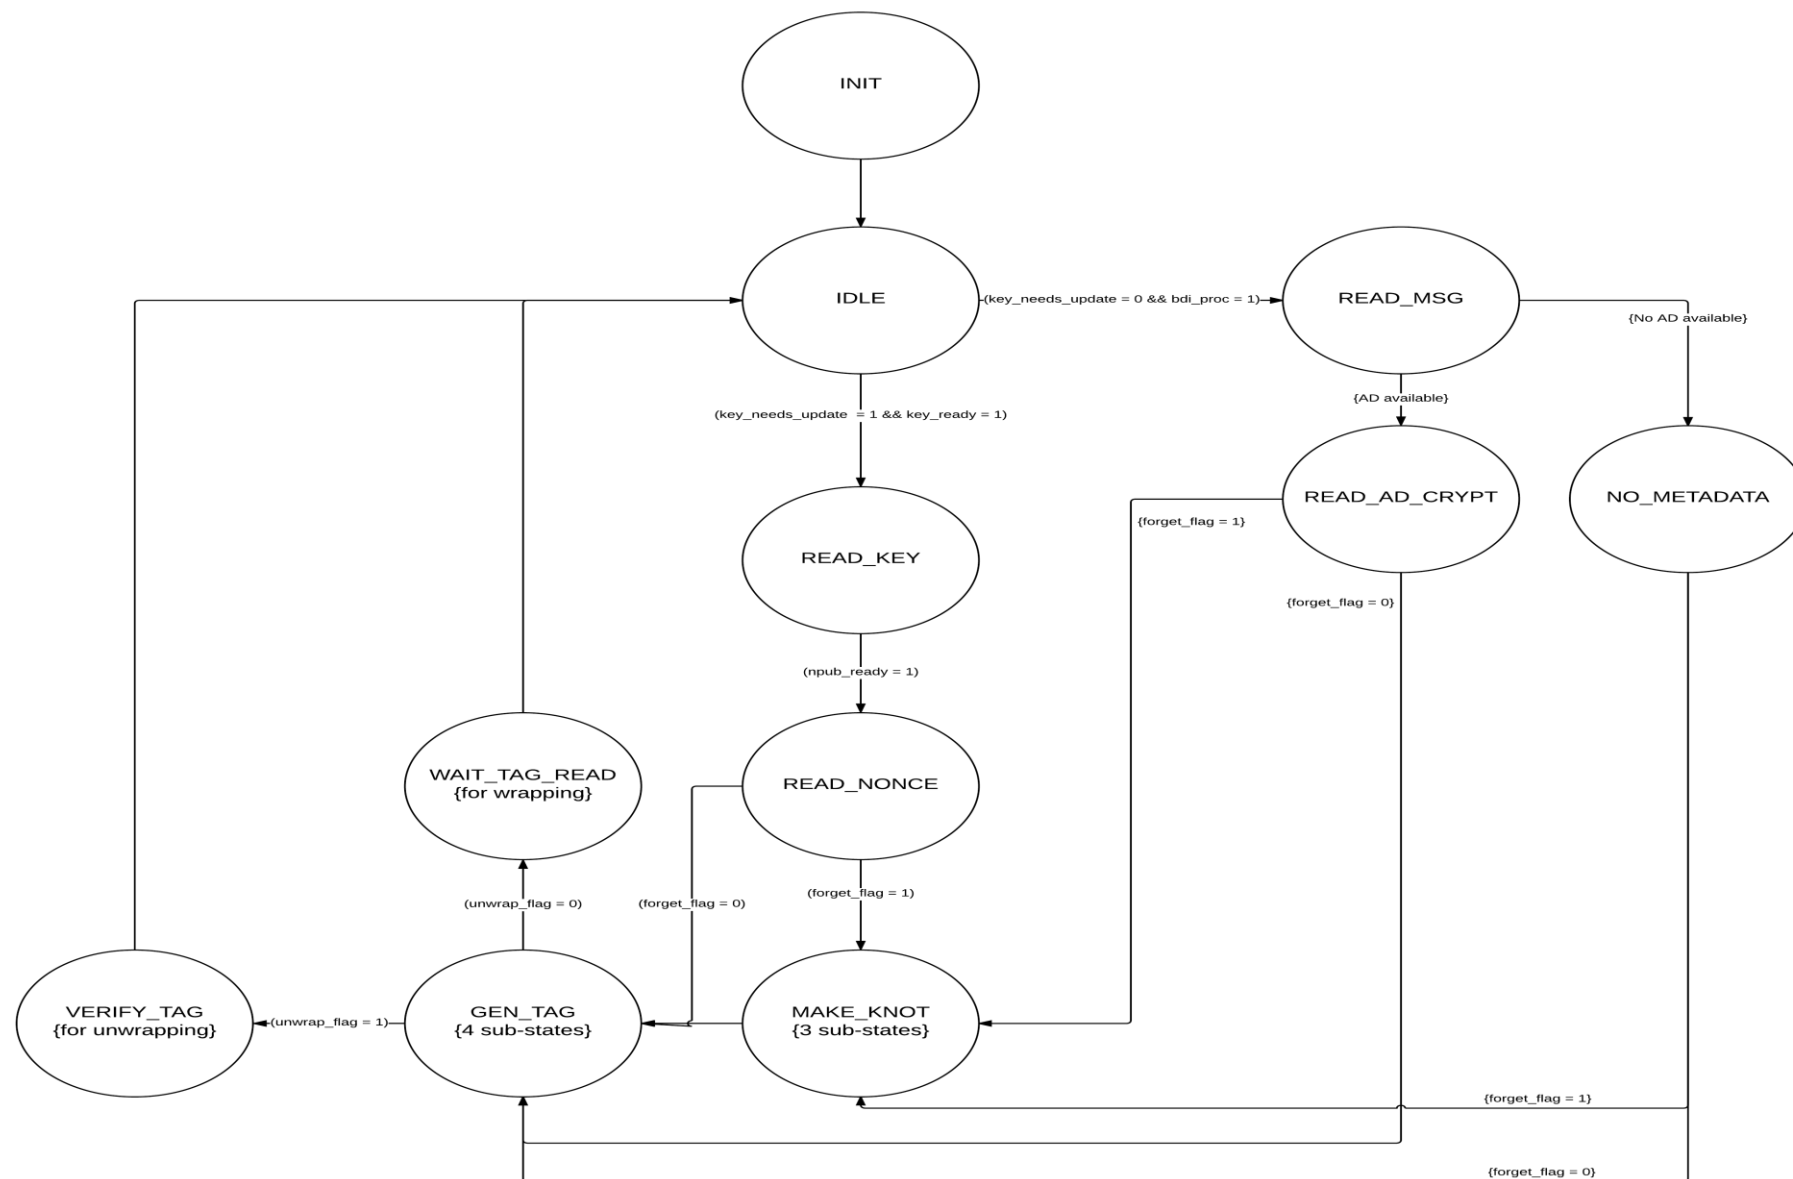
Design Decisions

- Input and Metadata limitations
 - GMU API specifies decrypted plaintext held in AUX FIFO until auth.
 - Hence can ciphertext per *Wrap* call limited to size of AUX FIFO
 - Decided to impose general inputdata limit:
 - Initially squeeze rate R_s
 - Later $(R_s - \tau)$ since otherwise extra permutation call would be needed for τ bytes
 - Two versions of each named instance:
 - Version A
 - Also limit on metadata ($R_a - R_s$), results in reduced area, only 1 permutation call per *Wrap* call (since total data consumed is $((R_s - \tau) + (R_a - R_s)) < R_a$)
 - Version B
 - No limit on metadata, only limit on inputdata

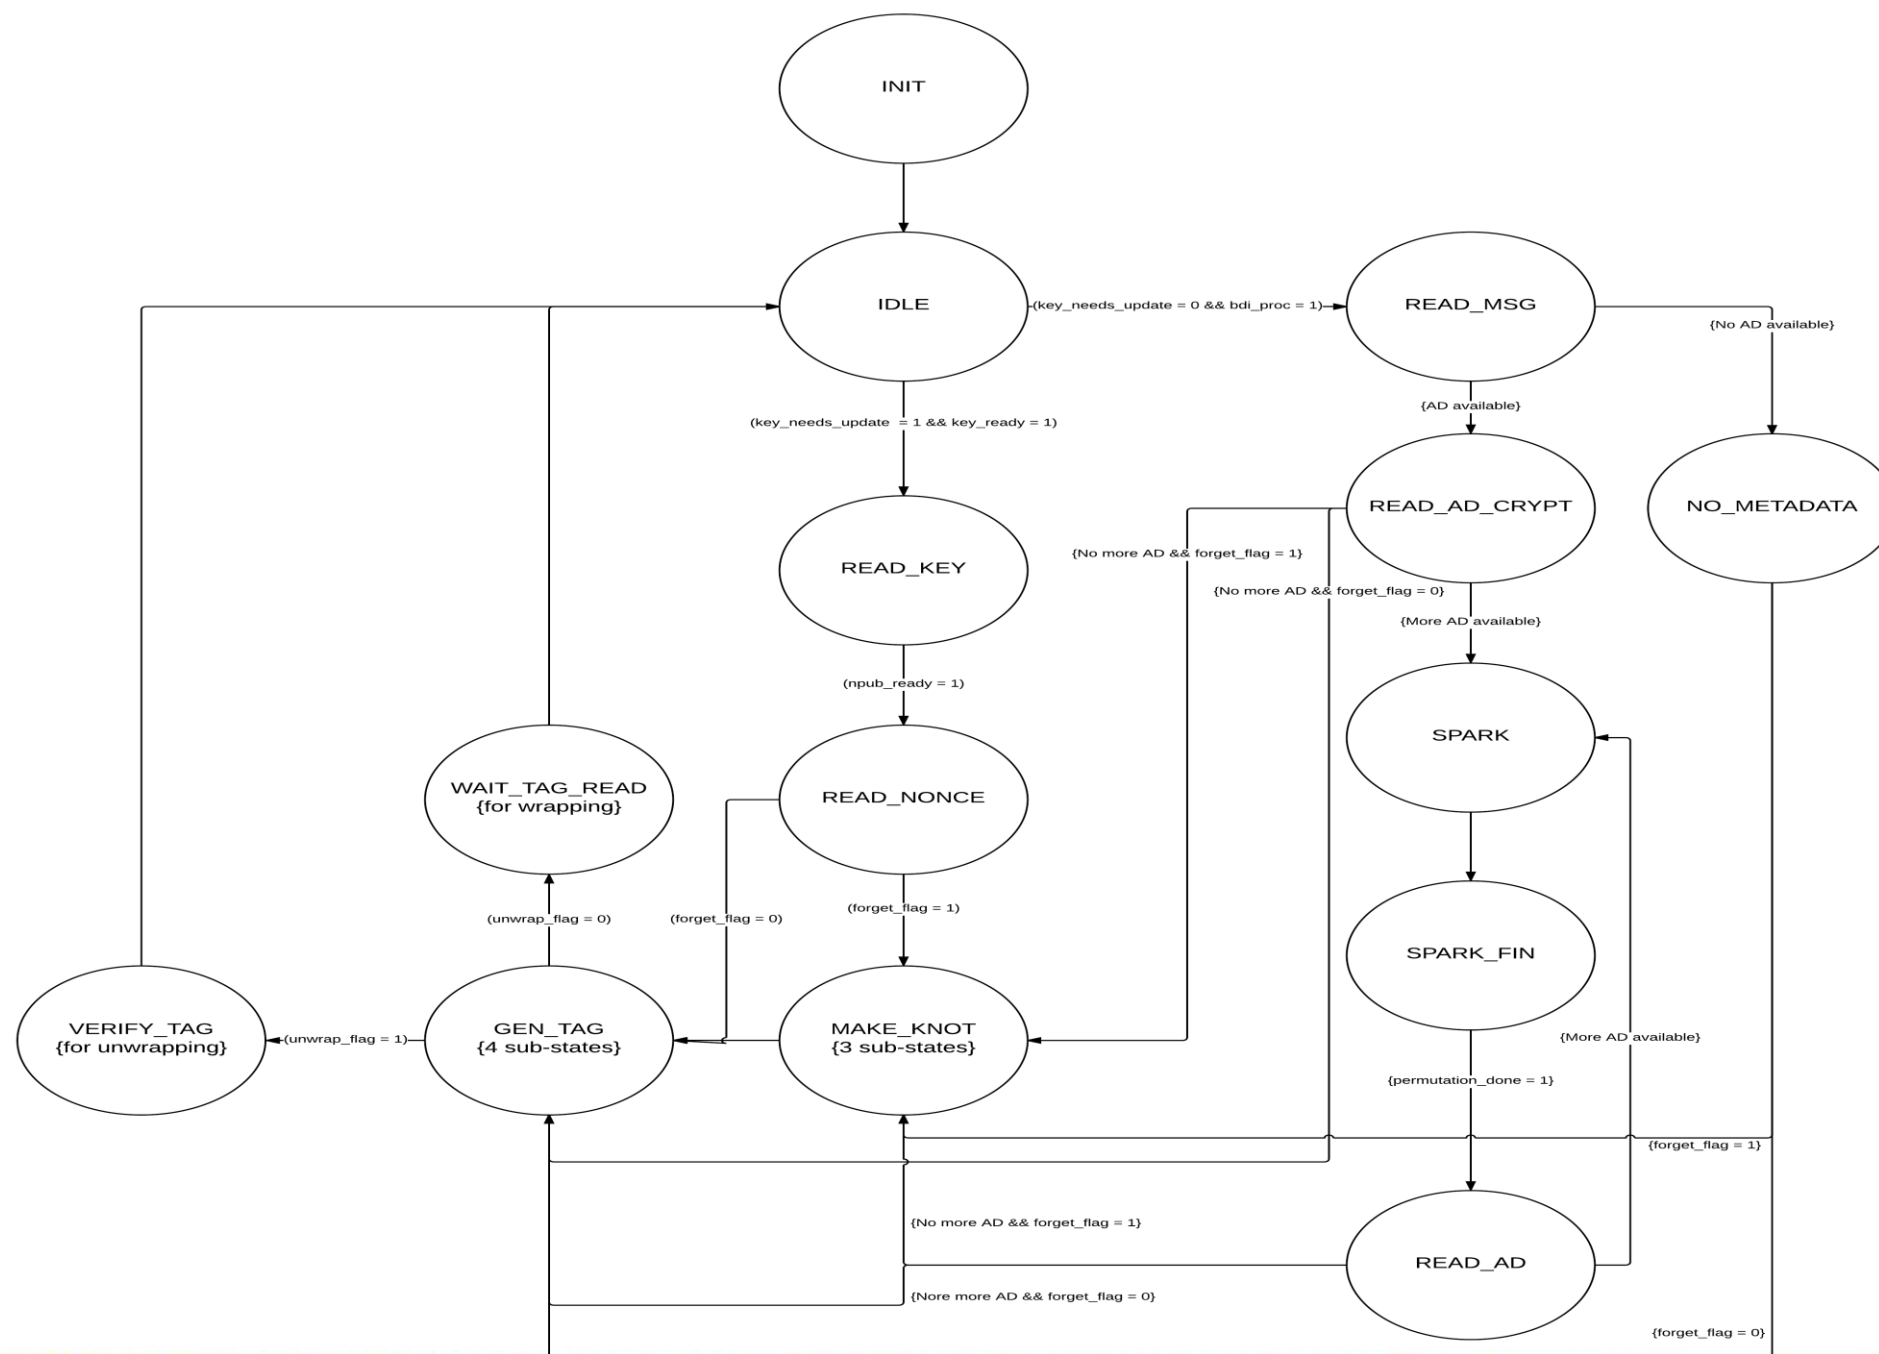
Design Decisions

- River & Lake Keyak have $\pi = 1$, hence no parallelism
 - Require logic for only 1 Piston
 - Eliminate related loops, certain functions (eg. *InjectCollective*)
 - Related arrays (eg. Et remembering how much output was used as tag or chaining value) reduced to single values
 - Static SUV diversification bytes (0x01 0x00)
 - Eliminate related conditional execution (eg. $(\pi > 1)$)
- Collapse sub-components into singular CipherCore
 - Piston, Engine, Motorist: Spread out over VHDL processes in CipherCore
 - Keccak_P round function as sub-component

Design Decisions – FSM (vA)



Design Decisions – FSM (vB)



Design Decisions

- State / Piston functionality in single VHDL process
 - Supports all required piston functionality (crypting, injection, etc.)
 - State is indexed in row/col/lane fashion, transformed to 8-bit byte-wise state indexing variable ω when necessary
- Permutation is done by dedicated VHDL process
 - Drives Keccak_P round sub-component (operates on state register)
 - Iterates for #rounds
 - Signals for 'activate' and 'done'

Design Decisions

- Version B problem: metadata internal buffer
 - When we consume more than (Ra-Rs) metadata we don't consume full blocks per permutation call
 - Need to store remnant data internally for usage during next call
 - Use Ra-sized internal buffer, indexed using exhaustive muxer
 - Alternative: shift register (not explored)
- Diversions from GMU API
 - Separate ports for input and metadata for ease of processing
 - Support for custom flags (eg. forget_flag)
 - Support alternate data processing order (first inputdata, then metadata)

Test Benches

- Test Benches test 1 base case and 3 edge cases, wrapping + unwrapping for each
 - Base Case
 - Data generated using *KeccakTools* test vector generation code
 - Maximum input & (Ra-Rs) metadata
 - Edge Case 1 (vB only)
 - Maximum input & long metadata
 - Edge Case 2
 - Single input & single metadata byte
 - Edge Case 3
 - Some input & no metadata

Performance Figures

- Implementation was *reference* implementation
- Emphasis on legibility / understandability to implementers rather than particular area/throughput optimizations
- Synthesis for FPGA and ASIC *very* intensive when writing non-optimized VHDL code
- Area figures obtained through separate synthesis of round function and preliminary synthesis of 1 instance
- Throughput figures expressed in bits per clock cycle

Performance Figures - Area

Architecture	Area (#slice LUTs)
Keccak _p Round Function (River Keyak)	1,899
Keccak _p Round Function (Lake Keyak)	2,956

Lake Keyak (version A)	Count
RAM (20x64-bit)	1
Multipliers (20x6-bit)	1
Adders/Subtractors	75 (of various bit-sizes)
Comparators	1443 (mostly 32-bit)
Multiplexers	828884 (mostly 1-bit 2-to-1)
XORs	7761
Registers (1216 bit)	1
Registers (1200-bit)	1
Registers (128-bit)	4
Registers (64-bit)	25
Registers (32-bit)	2
Registers (8-bit)	1
Registers (1 bit)	15
Counters	1
FSMs	2

Performance Figures - Throughput

Architecture	Initialization	Wrap
Keyak (version A)	$4 + i * (3 + \#rounds) + j * 2$	$3 + i * (3 + \#rounds) + j * 2 + m$
Keyak (version B)	$4 + i * (3 + \#rounds) + j * 2$	$3 + i * (3 + \#rounds) + j * 2 + m * l$

Architecture	Throughput	Throughput rate (Mbit/s)
<i>Keccak_p</i> Round Function (River Keyak)	$\frac{statesize}{T_{clk}}$	6004.563*
<i>Keccak_p</i> Round Function (Lake Keyak)	$\frac{statesize}{T_{clk}}$	11967.359*

Architecture	Inputdata, Metadata	Throughput rate (bits/clockcycle)
River Keyak (version A)	Inputdata: 52 bytes Metadata: 28 bytes	37.647
River Keyak (version B)	Inputdata: 52 bytes Metadata: 192	56.163
Lake Keyak (version A)	Inputdata: 152 bytes Metadata: 24 bytes	74.105
Lake Keyak (version B)	Inputdata: 152 bytes Metadata: 384	87.510